

# Time series forecasting

This notebook adapts the Tensorflow tutorial on [Time series forecasting](#) to data generated from a model for epidemic processes.

## Steps

1. Imports and setup
2. Load and prepare the generated data
3. Baseline forecasting
4. Univariate *GRU* based forecasting
5. Multivariate *GRU* based forecasting - Single Step
6. Multivariate *GRU* based forecasting - Multiple Steps

## Imports and setup

```
In [ ]: import tensorflow as tf

import matplotlib as mpl
import matplotlib.pyplot as plt
import numpy as np
import os
import pandas as pd

mpl.rcParams['figure.figsize'] = (8, 6)
mpl.rcParams['axes.grid'] = False
```

## Load and prepare the generated data

We load data from the ODE model introduced in the notebook "Probability and Information Theory". For each of the 150 virtual outbreaks (randomized and with different model parameters), we have time series (with 500 steps) for four the variables "Susceptible", "Infected", "Recovered", and "Deceased".

```
In [ ]: csv_path = "../data/epidemic_process_raw_data.csv"
df = pd.read_csv(csv_path)
df.head()
```

```
Out[ ]:
```

	1	2	3	4	5	6	7	
0	100.287149	103.541223	95.879814	96.354848	96.980932	97.855310	98.940537	100.18749
1	0.993774	1.017558	1.070030	1.116168	1.142078	1.134735	1.182418	1.27231
2	0.000000	0.017741	0.036585	0.054735	0.074266	0.096065	0.117691	0.13918
3	0.000000	0.000178	0.000364	0.000562	0.000757	0.000947	0.001160	0.00138
4	103.489688	100.282780	96.634270	98.532514	99.089272	97.440900	98.416534	101.40490

5 rows × 501 columns

```
In [ ]: dfSusceptible = df[df.index % 4 == 0]
dfSusceptible.head()
```

```
Out[ ]:
```

	1	2	3	4	5	6	7	
0	100.287149	103.541223	95.879814	96.354848	96.980932	97.855310	98.940537	100.18749
4	103.489688	100.282780	96.634270	98.532514	99.089272	97.440900	98.416534	101.40490
8	101.527421	97.711732	96.168179	95.677962	95.575326	96.109792	96.943831	98.000000
12	101.061107	99.112815	106.651686	101.622904	97.726686	95.692173	97.438263	102.000000
16	101.957189	101.898022	100.881113	99.892000	98.939878	98.048565	98.220024	99.200000

5 rows × 501 columns

```
In [ ]: dfInfected = df[df.index % 4 == 1]
dfInfected.head()
```

```
Out[ ]:
```

	1	2	3	4	5	6	7	8	
1	0.993774	1.017558	1.070030	1.116168	1.142078	1.134735	1.182418	1.272310	1.35656
5	1.021677	1.045410	1.120324	1.175914	1.236878	1.306676	1.387931	1.477973	1.54987
9	1.020043	1.011238	1.031122	1.048642	1.049479	1.022891	1.035862	1.079177	1.11453
13	1.035248	1.014189	1.133178	1.135622	1.157984	1.213088	1.281406	1.359858	1.42323
17	1.012666	1.016949	1.053194	1.097599	1.143640	1.192369	1.238880	1.283688	1.32430

5 rows × 501 columns

```
In [ ]: dfRecovered = df[df.index % 4 == 2]
dfRecovered.head()
```

```
Out[ ]:
```

	1	2	3	4	5	6	7	8	9	
2	0.0	0.017741	0.036585	0.054735	0.074266	0.096065	0.117691	0.139184	0.163615	0.1
6	0.0	0.017909	0.035748	0.056118	0.076620	0.097338	0.119592	0.143024	0.171253	0.2
10	0.0	0.016990	0.034644	0.052866	0.071444	0.090609	0.108733	0.126058	0.142408	0.1
14	0.0	0.017002	0.036315	0.057484	0.078381	0.098831	0.119563	0.140509	0.166486	0.1
18	0.0	0.017589	0.037434	0.056572	0.076275	0.096907	0.116533	0.135387	0.157592	0.1

5 rows × 501 columns

```
In [ ]: dfDead = df[df.index % 4 == 3]
dfDead.head()
```

```
Out[ ]:
```

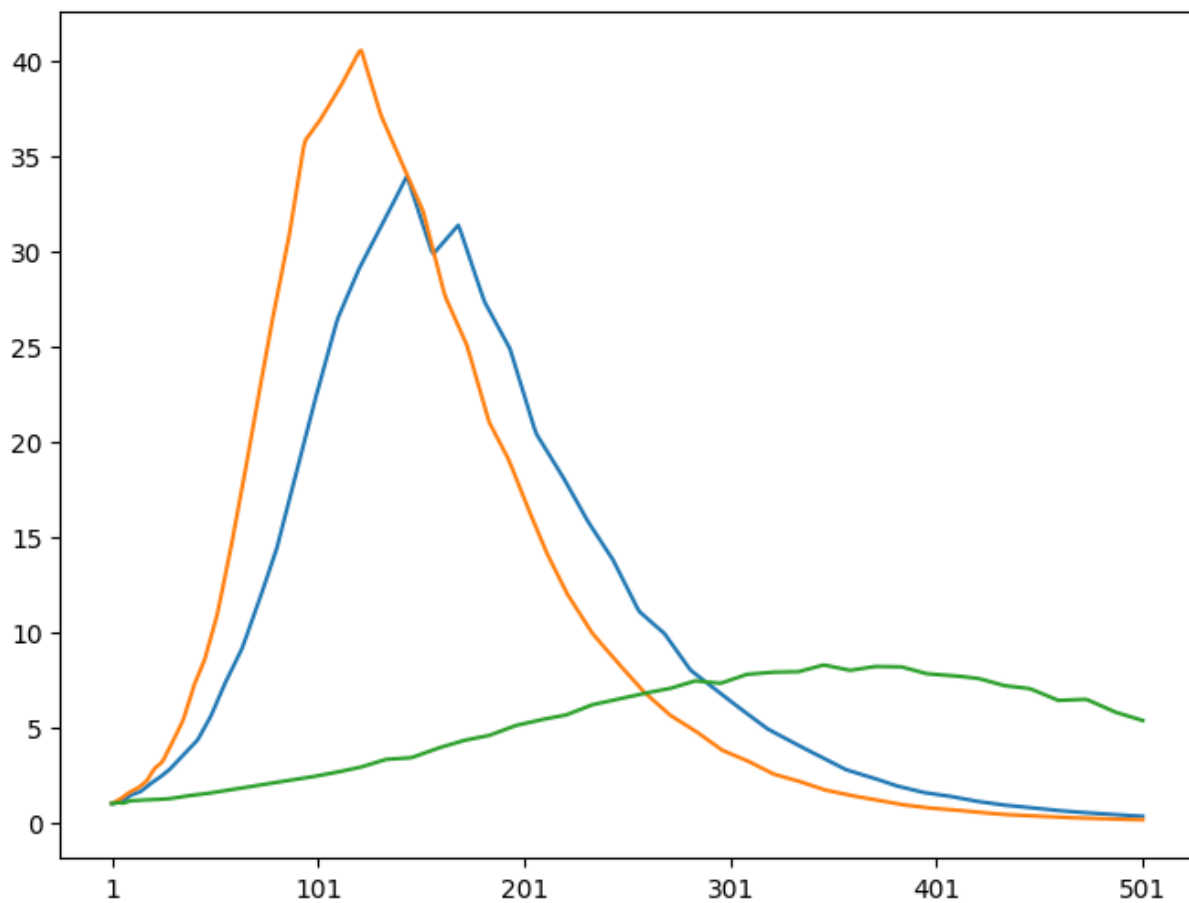
	1	2	3	4	5	6	7	8	9	
3	0.0	0.000178	0.000364	0.000562	0.000757	0.000947	0.001160	0.001389	0.001635	0.0
7	0.0	0.000175	0.000351	0.000558	0.000763	0.000968	0.001196	0.001443	0.001719	0.0
11	0.0	0.000171	0.000352	0.000538	0.000729	0.000927	0.001126	0.001324	0.001488	0.0
15	0.0	0.000181	0.000364	0.000563	0.000774	0.001003	0.001192	0.001351	0.001575	0.0
19	0.0	0.000180	0.000358	0.000550	0.000740	0.000931	0.001138	0.001359	0.001574	0.0

5 rows × 501 columns

Below a plot of three infection time series for the three first outbreaks.

```
In [ ]: dfInfected.loc[1,:].plot()
dfInfected.loc[5,:].plot()
dfInfected.loc[9,:].plot()
```

```
Out[ ]: <AxesSubplot:>
```



We define a 90% / 10% of data for training / testing.

```
In [ ]: dfInfected_arr = dfInfected.values
dfInfected_arr.shape
TRAIN_SPLIT = int(dfInfected_arr.shape[0]-dfInfected_arr.shape[0]*0.1)
TRAIN_SPLIT
```

Out[ ]: 135

We standardize the data.

```
In [ ]: uni_train_mean = dfInfected_arr[:TRAIN_SPLIT].mean()
uni_train_std = dfInfected_arr[:TRAIN_SPLIT].std()
uni_data = (dfInfected_arr-uni_train_mean)/uni_train_std
print ('\n Univariate data shape')
print(uni_data.shape)
```

Univariate data shape  
(150, 501)

We split the data into time series of `univariate_past_history=20` days length and predict the future of the current day, i.e., `univariate_future_target=0`, for the "infected" variable.

```
In [ ]: def univariate_data(dataset, start_series, end_series, history_size, target_size):
        data = []
        labels = []
        start_index = history_size
        end_index = len(dataset[0]) - target_size
        for c in range(start_series, end_series):
            for i in range(start_index, end_index):
                indices = range(i-history_size, i)
                # Reshape data from (history_size,) to (history_size, 1)
                data.append(np.reshape(dataset[c][indices], (history_size, 1)))
                labels.append(dataset[c][i+target_size])
        return np.array(data), np.array(labels)
```

```
In [ ]: univariate_past_history = 20 #days
        univariate_future_target = 0 #current day

        x_train_uni, y_train_uni = univariate_data(uni_data, 0, TRAIN_SPLIT,
                                                    univariate_past_history,
                                                    univariate_future_target)
        x_val_uni, y_val_uni = univariate_data(uni_data, TRAIN_SPLIT, len(uni_data),
                                                univariate_past_history,
                                                univariate_future_target)
```

```
In [ ]: print ('Single window of past history')
        print (x_train_uni[0])
        print ('\n Target number to predict')
        print (y_train_uni[0])
        print ('\n Number of traing data points')
        print (y_train_uni.shape[0])
        print ('\n Number of test data points')
        print (x_val_uni.shape[0])
```

Single window of past history

```
[[-0.95291296]
 [-0.95044298]
 [-0.94499366]
 [-0.9402021 ]
 [-0.93751136]
 [-0.93827393]
 [-0.93332191]
 [-0.92398652]
 [-0.91523643]
 [-0.90667772]
 [-0.90243571]
 [-0.89846308]
 [-0.89449045]
 [-0.89051782]
 [-0.88593997]
 [-0.87701137]
 [-0.86808277]
 [-0.85915417]
 [-0.85022557]
 [-0.84167481]]
```

Target number to predict

-0.8339932964893617

Number of traing data points

64935

Number of test data points

7215

*Create test data where the infection reached the peak*

```
In [ ]: peak_idx = np.argmax(uni_data, axis=1)
        peak_uni = []
        for i in range(len(uni_data)):
            idx = peak_idx[i]
            start = max(idx - univariate_past_history, 0)
            end = min(idx + univariate_past_history, uni_data.shape[1])
            temp_peak_uni = []
            for j in range(start, end+1):
                temp_peak_uni.append(uni_data[i,j])
            peak_uni.append(np.array(temp_peak_uni))

        peak_uni = np.array(peak_uni)
```

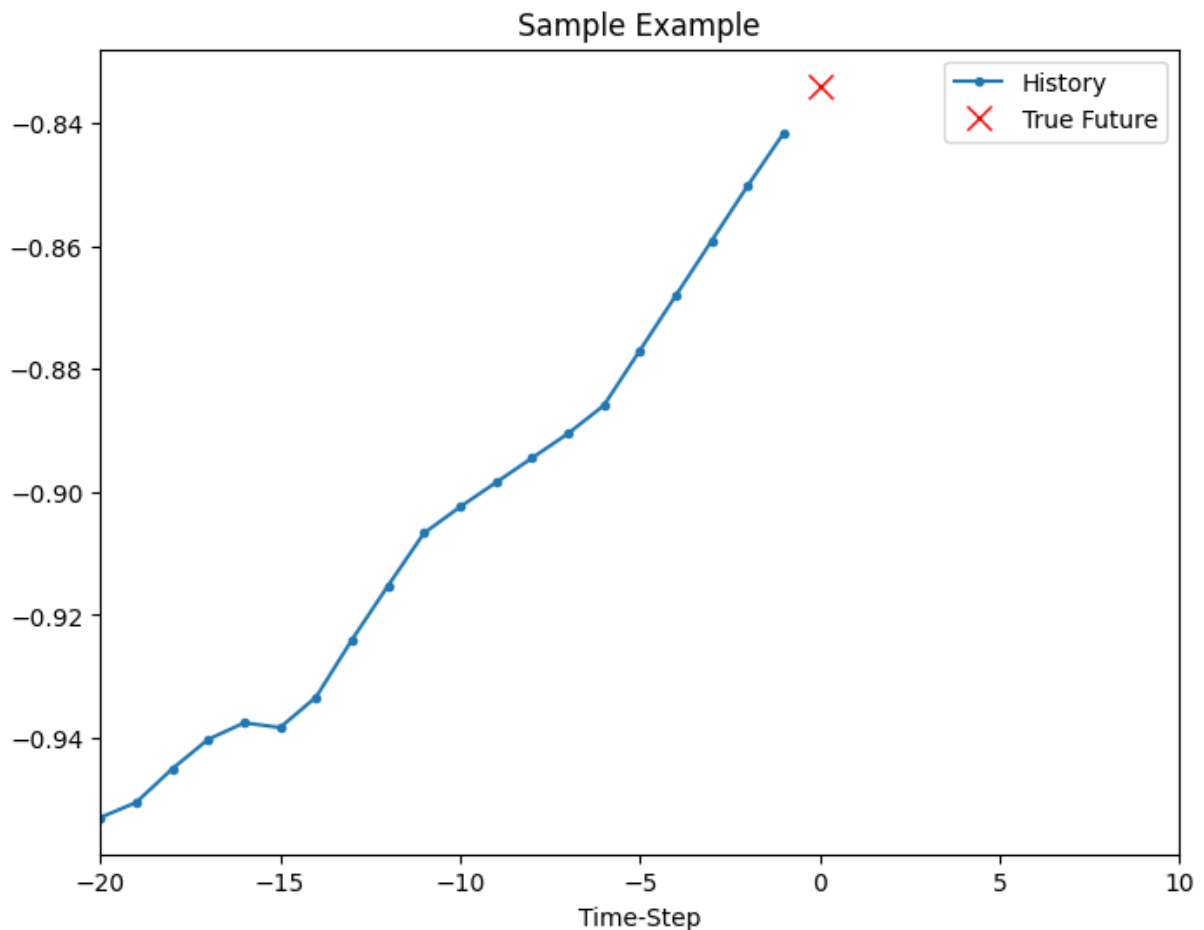
```
In [ ]: x_peak_uni, y_peak_uni = univariate_data(peak_uni, 0, len(uni_data),
                                                univariate_past_history,
                                                univariate_future_target)
```

```
In [ ]: def create_time_steps(length):
        return list(range(-length, 0))
```

```
In [ ]: def show_plot(plot_data, delta, title):
        labels = ['History', 'True Future', 'Model Prediction']
        marker = ['.-', 'rx', 'go']
        time_steps = create_time_steps(plot_data[0].shape[0])
        if delta:
            future = delta
        else:
            future = 0
        plt.title(title)
        for i, x in enumerate(plot_data):
            if i:
                plt.plot(future, plot_data[i], marker[i], markersize=10, label=labels[i])
            else:
                plt.plot(time_steps, plot_data[i].flatten(), marker[i], label=labels[i])
        plt.legend()
        plt.xlim([time_steps[0], (future+5)*2])
        plt.xlabel('Time-Step')
        return plt
```

```
In [ ]: show_plot([x_train_uni[0], y_train_uni[0]], 0, 'Sample Example')
```

```
Out[ ]: <module 'matplotlib.pyplot' from '/media/home/ngundermann/workspace/DML_notebooks/v
env/lib/python3.10/site-packages/matplotlib/pyplot.py'>
```



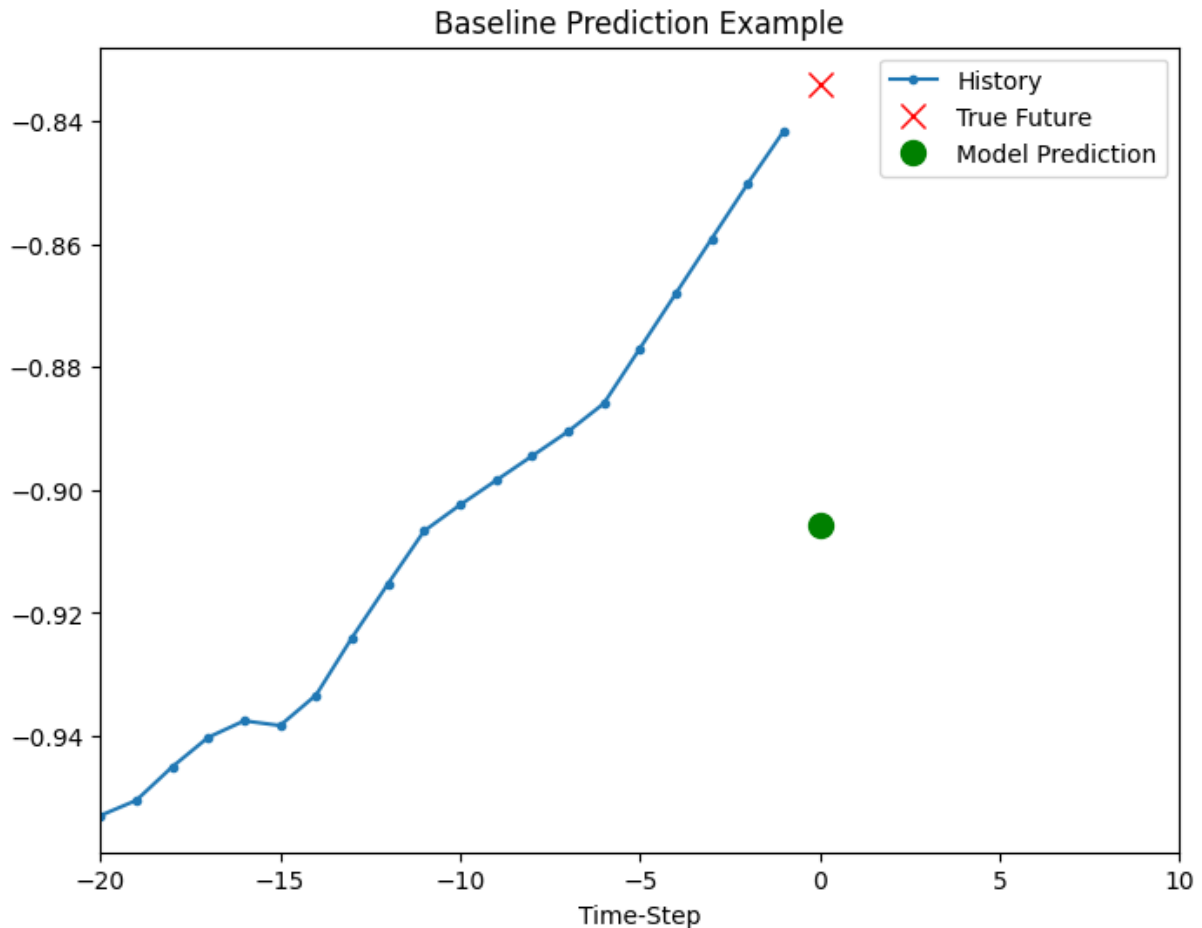
Baseline forecasting

Predicts the mean of the history .

```
In [ ]: def baseline(history):  
        return np.mean(history)
```

```
In [ ]: show_plot([x_train_uni[0], y_train_uni[0], baseline(x_train_uni[0])], 0, 'Baseline
```

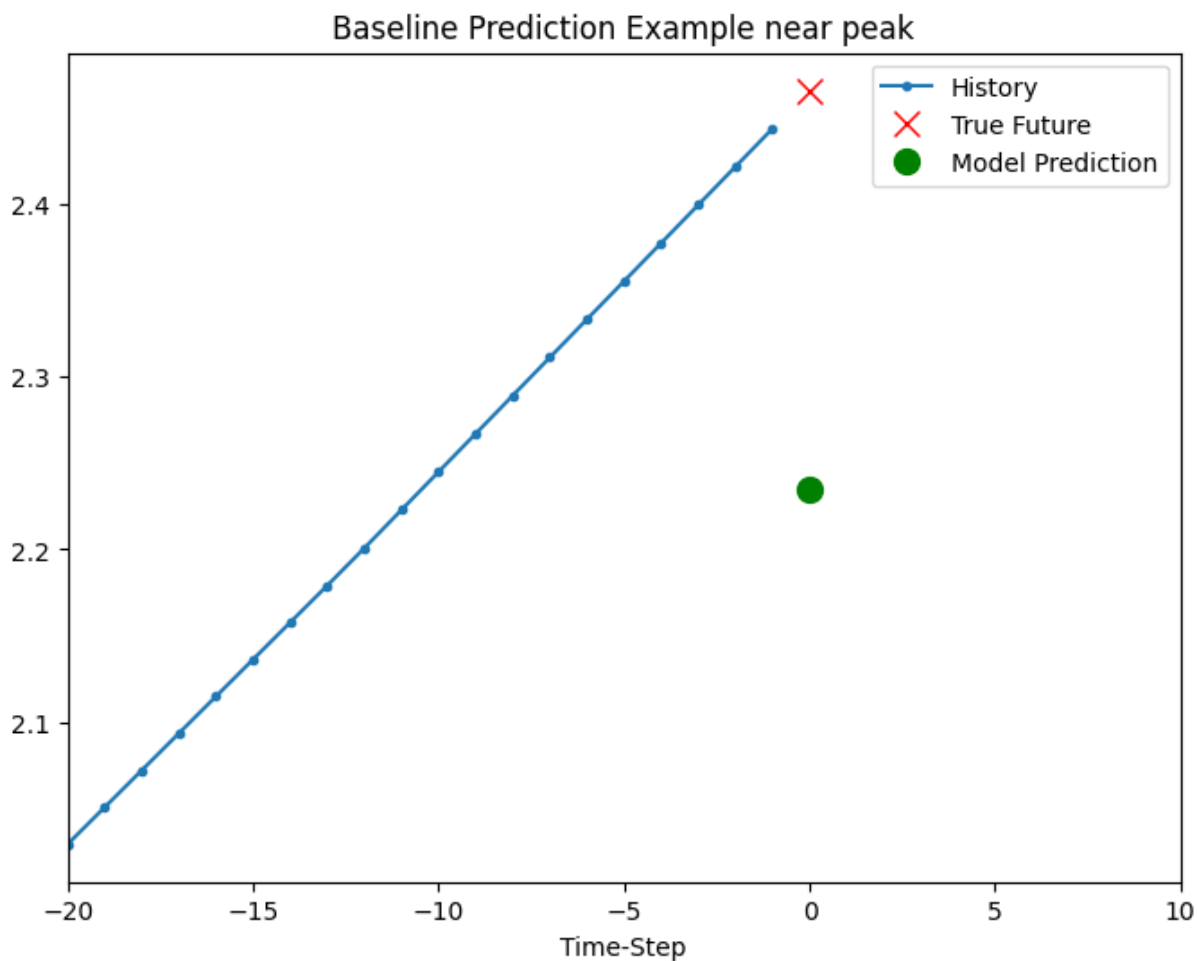
```
Out[ ]: <module 'matplotlib.pyplot' from '/media/home/ngundermann/workspace/DML_notebooks/v  
env/lib/python3.10/site-packages/matplotlib/pyplot.py'>
```



```
In [ ]: show_plot([x_peak_uni[0], y_peak_uni[0], baseline(x_peak_uni[0])], 0, 'Baseline Pre
```

```
Out[ ]: <module 'matplotlib.pyplot' from '/media/home/ngundermann/workspace/DML_notebooks/v  
env/lib/python3.10/site-packages/matplotlib/pyplot.py'>
```





## Univariate *GRU* based forecasting

```
In [ ]: print (x_train_uni.shape)
        print (y_train_uni.shape)
        x_train_uni.dtype
```

```
(64935, 20, 1)
```

```
(64935,)
```

```
Out[ ]: dtype('float64')
```

Batching and resampling; the dataset is repeated indefinitely. Check the tutorial for the details.

```
In [ ]: BATCH_SIZE = 256
        BUFFER_SIZE = 10000

        train_univariate = tf.data.Dataset.from_tensor_slices((x_train_uni, y_train_uni))
        train_univariate = train_univariate.cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE).

        val_univariate = tf.data.Dataset.from_tensor_slices((x_val_uni, y_val_uni))
        val_univariate = val_univariate.batch(BATCH_SIZE).repeat()

        peak_val_univariate = tf.data.Dataset.from_tensor_slices((x_peak_uni, y_peak_uni))
        peak_val_univariate = peak_val_univariate.batch(1).repeat()

        train_univariate
```

```
Out[ ]: <_RepeatDataset element_spec=(TensorSpec(shape=(None, 20, 1), dtype=tf.float64, nam
e=None), TensorSpec(shape=(None,), dtype=tf.float64, name=None))>
```

We define the first *GRU* model with 8 units.

```
In [ ]: simple_lstm_model = tf.keras.models.Sequential([
        tf.keras.layers.GRU(8, input_shape=x_train_uni.shape[-2:]),
        tf.keras.layers.Dense(1)
    ])

    simple_lstm_model.compile(optimizer='adam', loss='mae')
    simple_lstm_model.summary()
    x_train_uni.shape[-2:]
```

/media/home/ngundermann/workspace/DML\_notebooks/venv/lib/python3.10/site-packages/ke  
ras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input\_shape`/`input\_dim`  
argument to a layer. When using Sequential models, prefer using an `Input(shape)` ob  
ject as the first layer in the model instead.

```
    super().__init__(**kwargs)
```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
gru_12 (GRU)	(None, 8)	264
dense_9 (Dense)	(None, 1)	9

Total params: 273 (1.07 KB)

Trainable params: 273 (1.07 KB)

Non-trainable params: 0 (0.00 B)

```
Out[ ]: (20, 1)
```

*The GRU layer comes with less parameters than the LSTM layer (320) we have seen in the original notebook. This is because the GRU architecture is much simpler than LSTMs'.*

```
In [ ]: for x, y in val_univariate.take(1):
        print(simple_lstm_model.predict(x).shape)
        print(y.shape)
```

8/8 ————— 0s 1ms/step

(256, 1)

(256,)

2024-05-30 09:00:21.924564: W tensorflow/core/framework/local\_rendevvous.cc:404] Local rendezvous is aborting with status: OUT\_OF\_RANGE: End of sequence

When passing an indefinitely repeated training data set, we need to specify the number of steps per training interval (epoch).

```
In [ ]: EVALUATION_INTERVAL = 2000
        EPOCHS = 10

        simple_lstm_model.fit(train_univariate,
                               epochs=EPOCHS,
                               steps_per_epoch=EVALUATION_INTERVAL,
                               validation_data=val_univariate,
                               validation_steps=50)
```

Epoch 1/10

2000/2000 ————— 8s 3ms/step - loss: 0.0792 - val\_loss: 0.0031

Epoch 2/10

2000/2000 ————— 7s 3ms/step - loss: 0.0029 - val\_loss: 0.0026

Epoch 3/10

2000/2000 ————— 7s 3ms/step - loss: 0.0024 - val\_loss: 0.0016

Epoch 4/10

2000/2000 ————— 7s 3ms/step - loss: 0.0022 - val\_loss: 0.0014

Epoch 5/10

2000/2000 ————— 7s 3ms/step - loss: 0.0020 - val\_loss: 0.0015

Epoch 6/10

2000/2000 ————— 7s 3ms/step - loss: 0.0019 - val\_loss: 0.0013

Epoch 7/10

2000/2000 ————— 7s 3ms/step - loss: 0.0017 - val\_loss: 0.0012

Epoch 8/10

2000/2000 ————— 7s 3ms/step - loss: 0.0017 - val\_loss: 0.0018

Epoch 9/10

2000/2000 ————— 6s 3ms/step - loss: 0.0017 - val\_loss: 8.7684e-04

Epoch 10/10

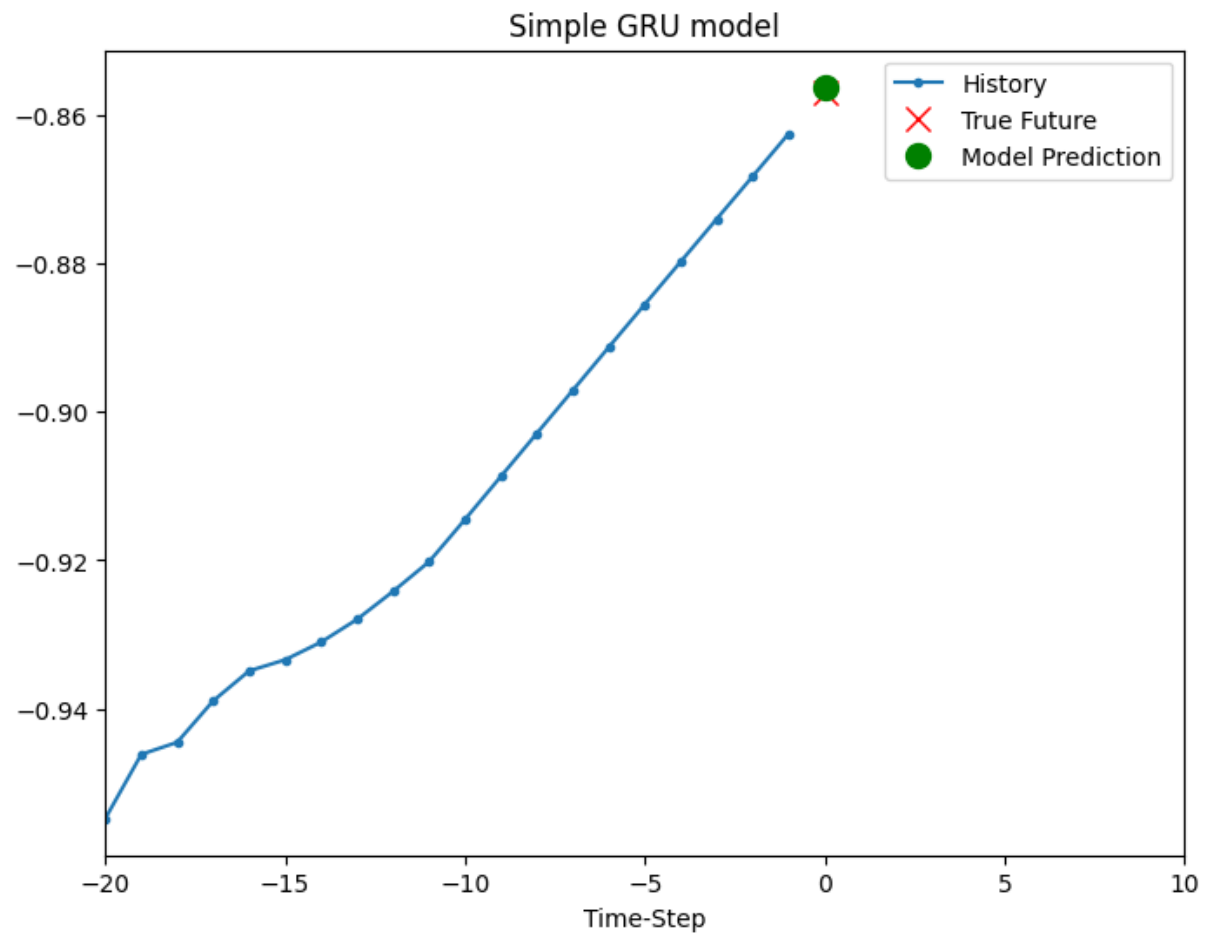
2000/2000 ————— 6s 3ms/step - loss: 0.0017 - val\_loss: 0.0012

Out[ ]: <keras.src.callbacks.history.History at 0x7f04384e0e50>

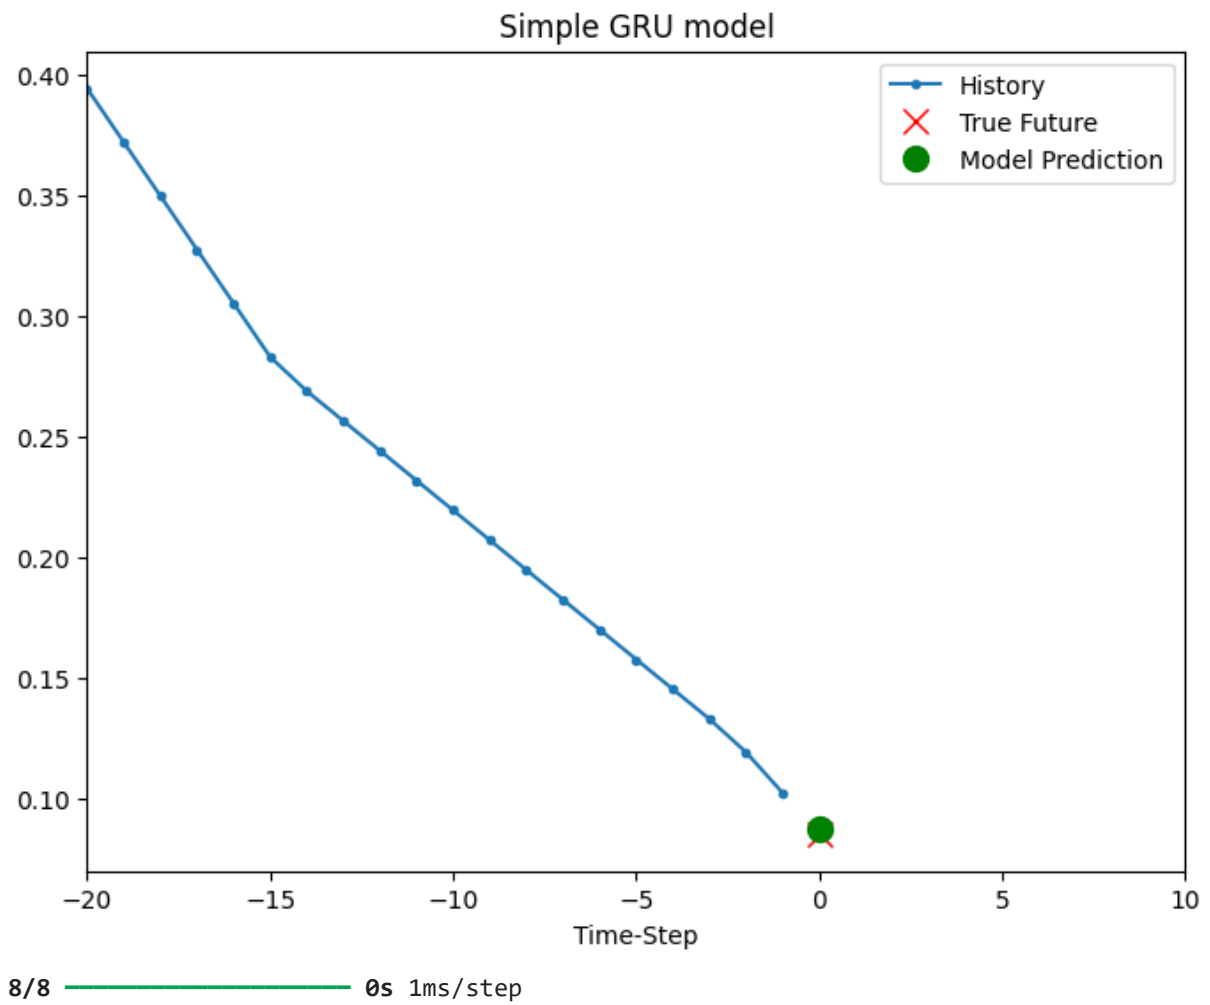
*The training loss went down very fast and converged after a few epochs. However, the validation loss started very low at the beginning, but did decrease quite fast. The training of one epoch finished faster, which could be caused by the fewer parameters of the model but also by the different machine on that the training was ran.*

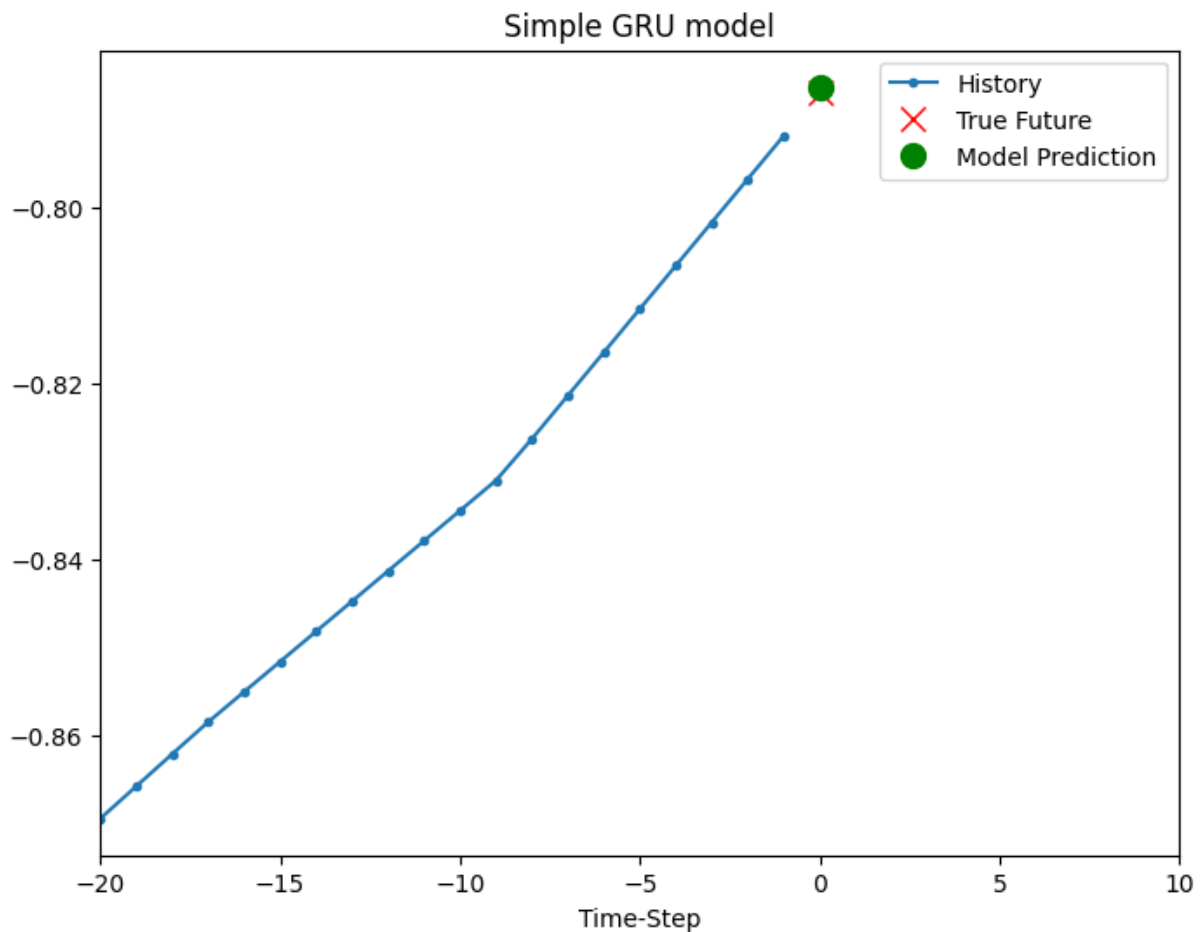
```
In [ ]: for x, y in val_univariate.take(3):
        plot = show_plot([x[0].numpy(), y[0].numpy(), simple_lstm_model.predict(x)[0]],
                          plot.show())
```

8/8 ————— 0s 1ms/step



8/8 — 0s 1ms/step





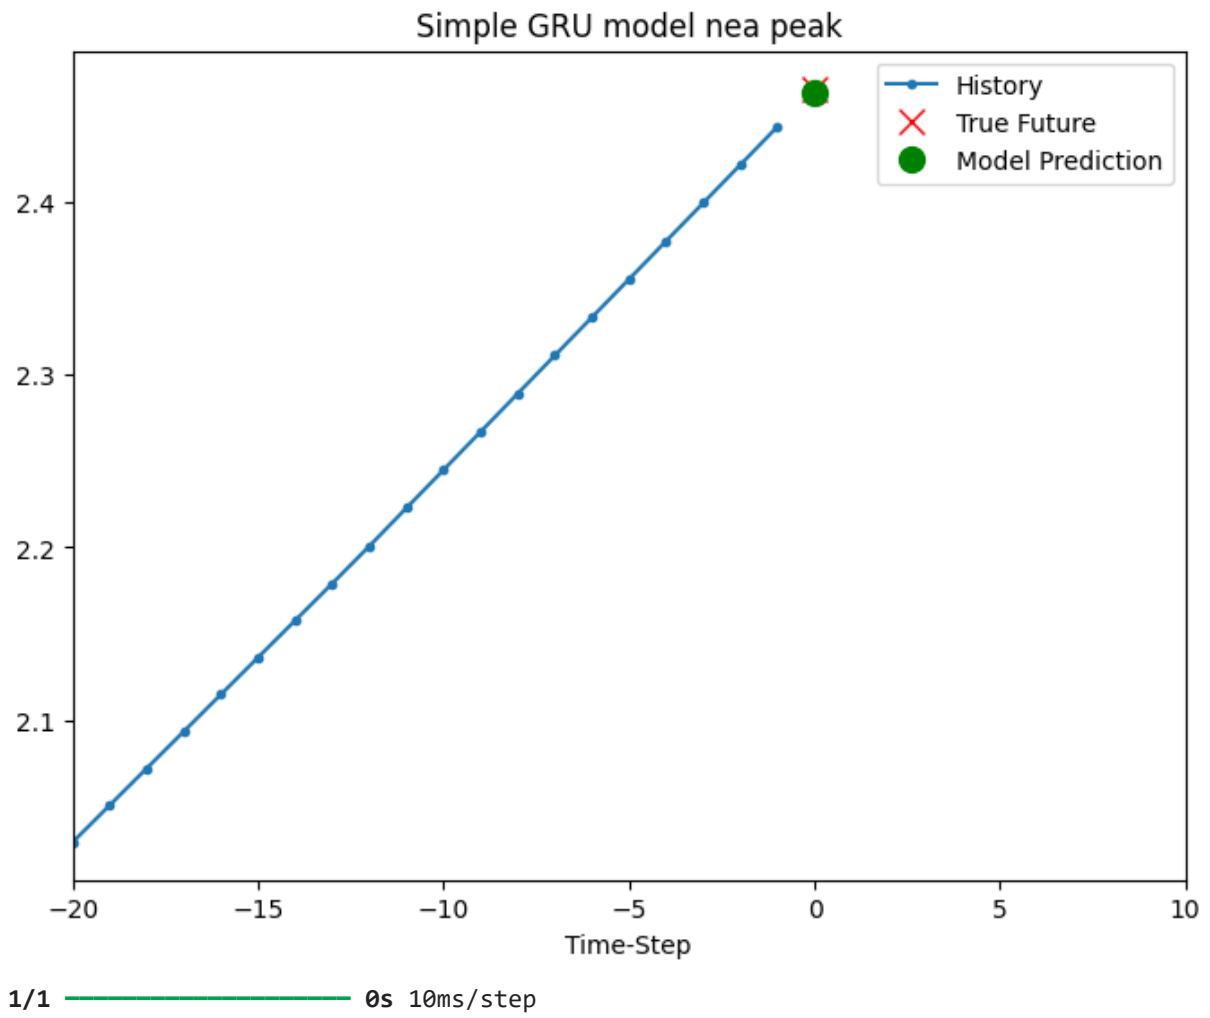
2024-05-30 09:01:29.437318: W tensorflow/core/framework/local\_rendevous.cc:404] Local rendezvous is aborting with status: OUT\_OF\_RANGE: End of sequence

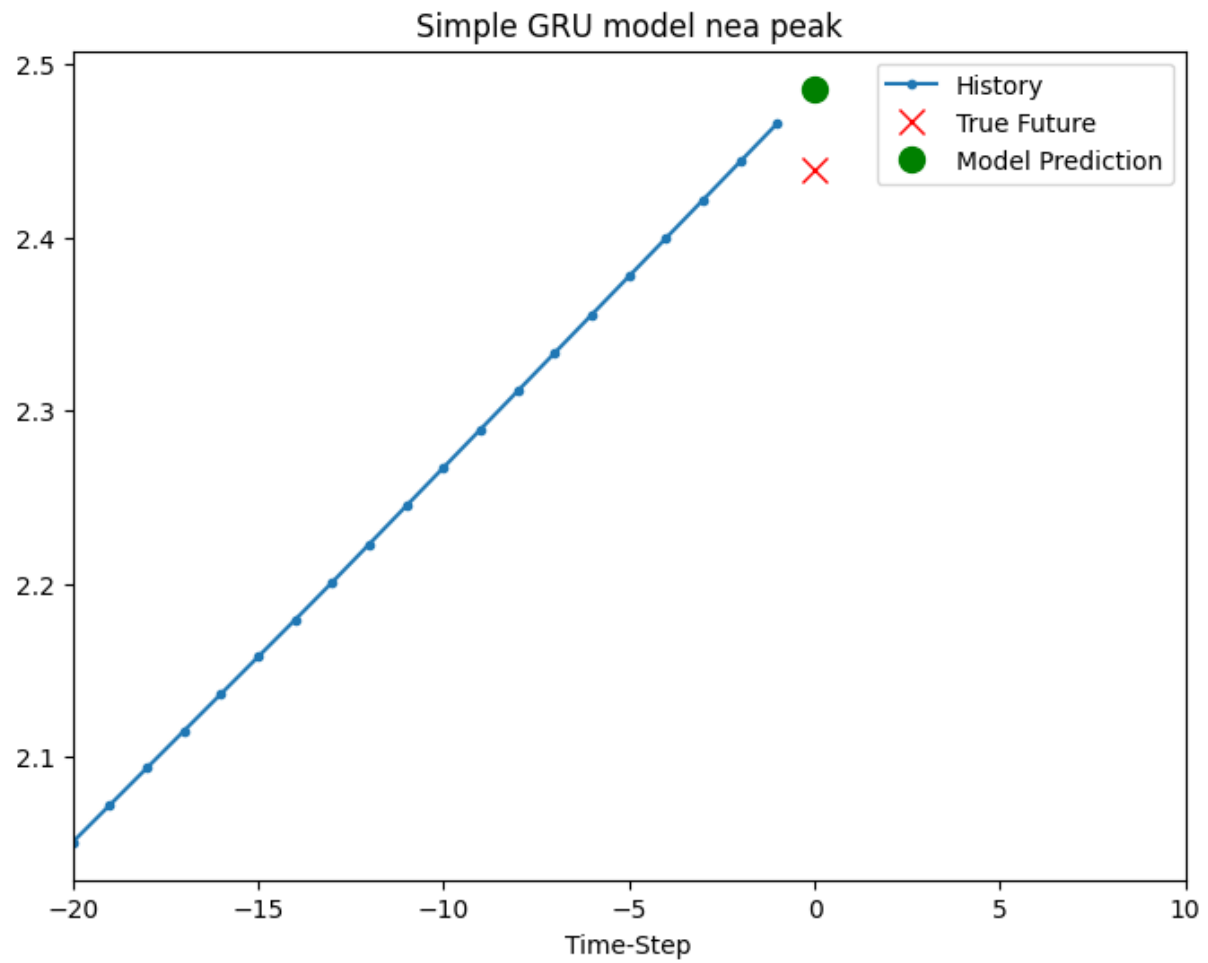
*The predictions look quite good, since the validation loss was quite low the whole time. The prediction are not exactly at the true value, but very close to it.*

*Diagrams with the peek of infection*

```
In [ ]: for x, y in peak_val_univariate.take(3):  
        plot = show_plot([x[0].numpy(), y[0].numpy(), simple_lstm_model.predict(x)[0]],  
                          plot.show())
```

1/1 ————— 0s 94ms/step





1/1 — 0s 11ms/step





2024-05-30 09:01:29.878153: W tensorflow/core/framework/local\_rendezvous.cc:404] Local rendezvous is aborting with status: OUT\_OF\_RANGE: End of sequence

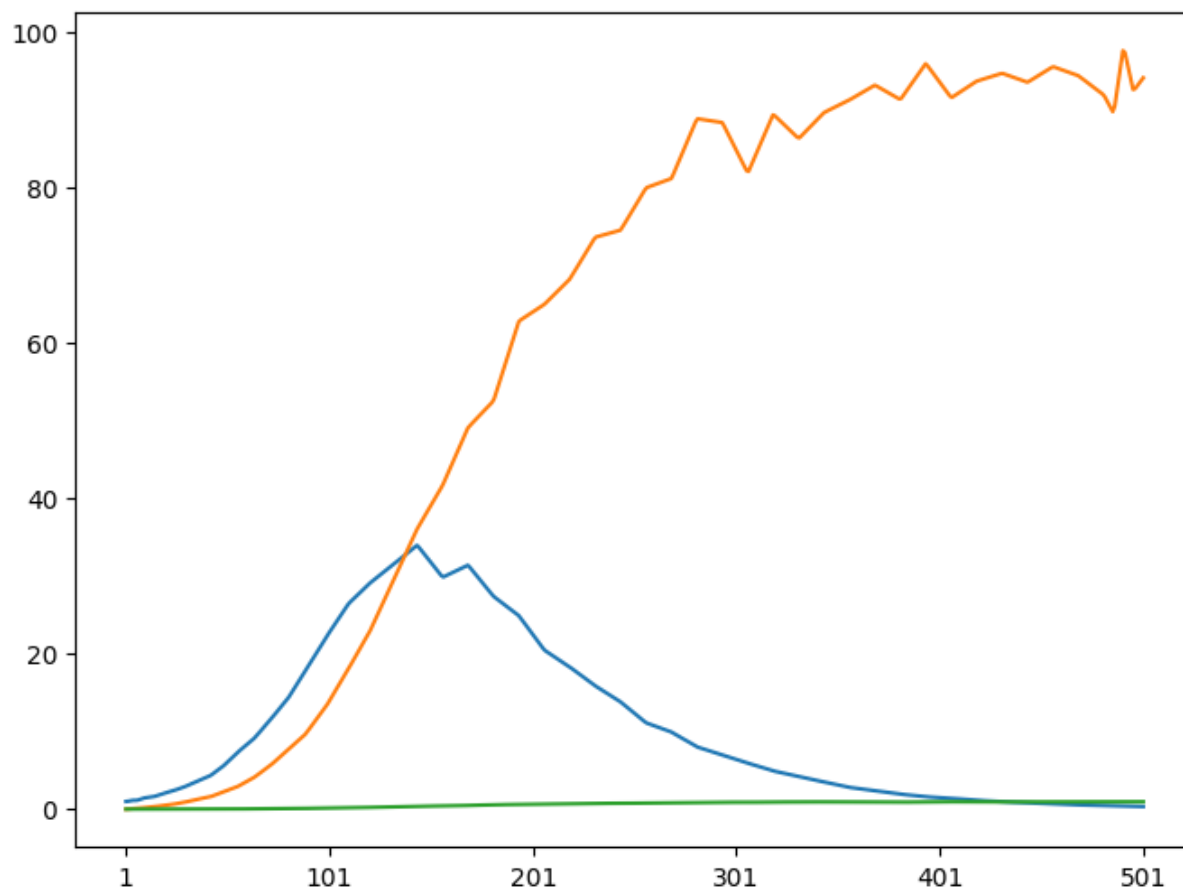
*Even the prediction near infection peak were not that good. It seems that the model has problems with changing slopes of the ground truths trajectory, i.e. when the infection starts to decline.*

## Multivariate *GRU* based forecasting - Single Step

We use three variables "Infected", "Recovered", and "Deceased", to forecast "Infected" at one single day in the future.

Here a plot of the time series of the three variables for one outbreak.

```
In [ ]: dfInfected.loc[1,:].plot()
dfRecovered.loc[2,:].plot()
dfDead.loc[3,:].plot()
dfInfected = dfInfected.values
dfRecovered_arr = dfRecovered.values
dfDead_arr = dfDead.values
```



We prepare the dataset.

```
In [ ]: #as before
dfInfected_train_mean = dfInfected_arr[:TRAIN_SPLIT].mean()
dfInfected_train_std = dfInfected_arr[:TRAIN_SPLIT].std()
dfInfected_data = (dfInfected_arr - dfInfected_train_mean) / dfInfected_train_std
#for Recovered
dfRecovered_train_mean = dfRecovered_arr[:TRAIN_SPLIT].mean()
dfRecovered_train_std = dfRecovered_arr[:TRAIN_SPLIT].std()
dfRecovered_data = (dfRecovered_arr - dfRecovered_train_mean) / dfRecovered_train_std
#for Dead
dfDead_train_mean = dfDead_arr[:TRAIN_SPLIT].mean()
dfDead_train_std = dfDead_arr[:TRAIN_SPLIT].std()
dfDead_data = (dfDead_arr - dfDead_train_mean) / dfDead_train_std
```

```
In [ ]: dataset = np.array([dfInfected_data, dfRecovered_data, dfDead_data])
dataset.shape
print('\n Multivariate data shape')
print(dataset.shape)
```

```
Multivariate data shape
(3, 150, 501)
```

```
In [ ]: def multivariate_data(dataset, target, start_series, end_series, history_size,
                                target_size, step, single_step=False):
    data = []
    labels = []
    start_index = history_size
    end_index = len(dataset[0][0]) - target_size
    for c in range(start_series, end_series):
        for i in range(start_index, end_index):
            indices = range(i-history_size, i, step)
            one = dataset[0][c][indices]
            two = dataset[1][c][indices]
            three = dataset[2][c][indices]
            data.append(np.transpose(np.array([one, two, three])))

            if single_step:
                labels.append(target[c][i+target_size])
            else:
                labels.append(np.transpose(target[c][i:i+target_size]))
    return np.array(data), np.array(labels)
```

We get training and validation data for time series with a `past_history = 20` days for every other day (`STEP = 2`) and want to predict the "Infected" five days ahead (`future_target = 5`).

```
In [ ]: past_history = 20
        future_target = 5
        STEP = 2

        x_train_single, y_train_single = multivariate_data(dataset, dfInfected_data, 0, TRAIN_SIZE,
                                                            past_history, future_target, STEP,
                                                            single_step=True)
        x_val_single, y_val_single = multivariate_data(dataset, dfInfected_data, TRAIN_SIZE,
                                                         past_history, future_target, STEP,
                                                         single_step=True)

In [ ]: print ('Single window of past history : {}'.format(x_train_single[0].shape))
        print(dataset.shape)
```

```
Single window of past history : (10, 3)
(3, 150, 501)
```

*As before, creating a dataset around the infection peak.*

```
In [ ]: def multivariate_data_near_peak_infection(dataset, past_history, future_target):
    peak_idx = np.argmax(dataset[0], axis=1)
    peak_infected = []
    peak_infected_recovered = []
    peak_infected_dead = []
    for i in range(len(dataset[1])):
        idx = peak_idx[i]
        start = max(idx - 1 - past_history - future_target // 2, 0)
        end = min(idx + future_target // 2 + future_target % 2, dataset.shape[2])
        temp_peak_infected = []
        temp_peak_infected_recovered = []
        temp_peak_infected_dead = []
        for j in range(start, end):
            temp_peak_infected.append(dataset[0,i,j])
            temp_peak_infected_recovered.append(dataset[1,i,j])
            temp_peak_infected_dead.append(dataset[2,i,j])
        peak_infected.append(np.array(temp_peak_infected))
        peak_infected_recovered.append(np.array(temp_peak_infected_recovered))
        peak_infected_dead.append(np.array(temp_peak_infected_dead))

    return np.array([peak_infected, peak_infected_recovered, peak_infected_dead])
```

```
In [ ]: peak_data = multivariate_data_near_peak_infection(dataset, past_history, future_target)
peak_data.shape
```

```
Out[ ]: (3, 150, 26)
```

```
In [ ]: x_peak_single, y_peak_single = multivariate_data(peak_data, peak_data[0], 0, peak_data[1],
    past_history, future_target, STEP,
    single_step=True)
```

```
In [ ]: x_peak_single.shape
```

```
Out[ ]: (150, 10, 3)
```

As before, batching and resampling; the dataset is repeated indefinitely.

```
In [ ]: train_data_single = tf.data.Dataset.from_tensor_slices((x_train_single, y_train_single))
train_data_single = train_data_single.cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)

val_data_single = tf.data.Dataset.from_tensor_slices((x_val_single, y_val_single))
val_data_single = val_data_single.batch(BATCH_SIZE).repeat()

peak_data_single = tf.data.Dataset.from_tensor_slices((x_peak_single, y_peak_single))
peak_data_single = peak_data_single.batch(1).repeat()
```

```
In [ ]: single_step_model = tf.keras.models.Sequential()
single_step_model.add(tf.keras.layers.GRU(32, input_shape=x_train_single.shape[-2:]
single_step_model.add(tf.keras.layers.Dense(1))

single_step_model.compile(optimizer=tf.keras.optimizers.RMSprop(), loss='mae')
single_step_model.summary()
x_train_single.shape[-2:]
```

```
/media/home/ngundermann/workspace/DML_notebooks/venv/lib/python3.10/site-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
```

```
super().__init__(**kwargs)
```

```
Model: "sequential_10"
```

Layer (type)	Output Shape	Param #
gru_13 (GRU)	(None, 32)	3,552
dense_10 (Dense)	(None, 1)	33

Total params: 3,585 (14.00 KB)

Trainable params: 3,585 (14.00 KB)

Non-trainable params: 0 (0.00 B)

```
Out[ ]: (10, 3)
```

*Again we got a network with fewer parameters, since the GRU architecture comes with less parameters than the LSTM architecture.*

```
In [ ]: for x, y in val_data_single.take(1):
        print(single_step_model.predict(x).shape)
        print('\n Number of traing data points')
        print(x_train_single.shape[0])
        print('\n Number of test data points')
        print(x_val_single.shape[0])
```

8/8 ————— 0s 1ms/step

(256, 1)

Number of traing data points  
64260

Number of test data points  
7140

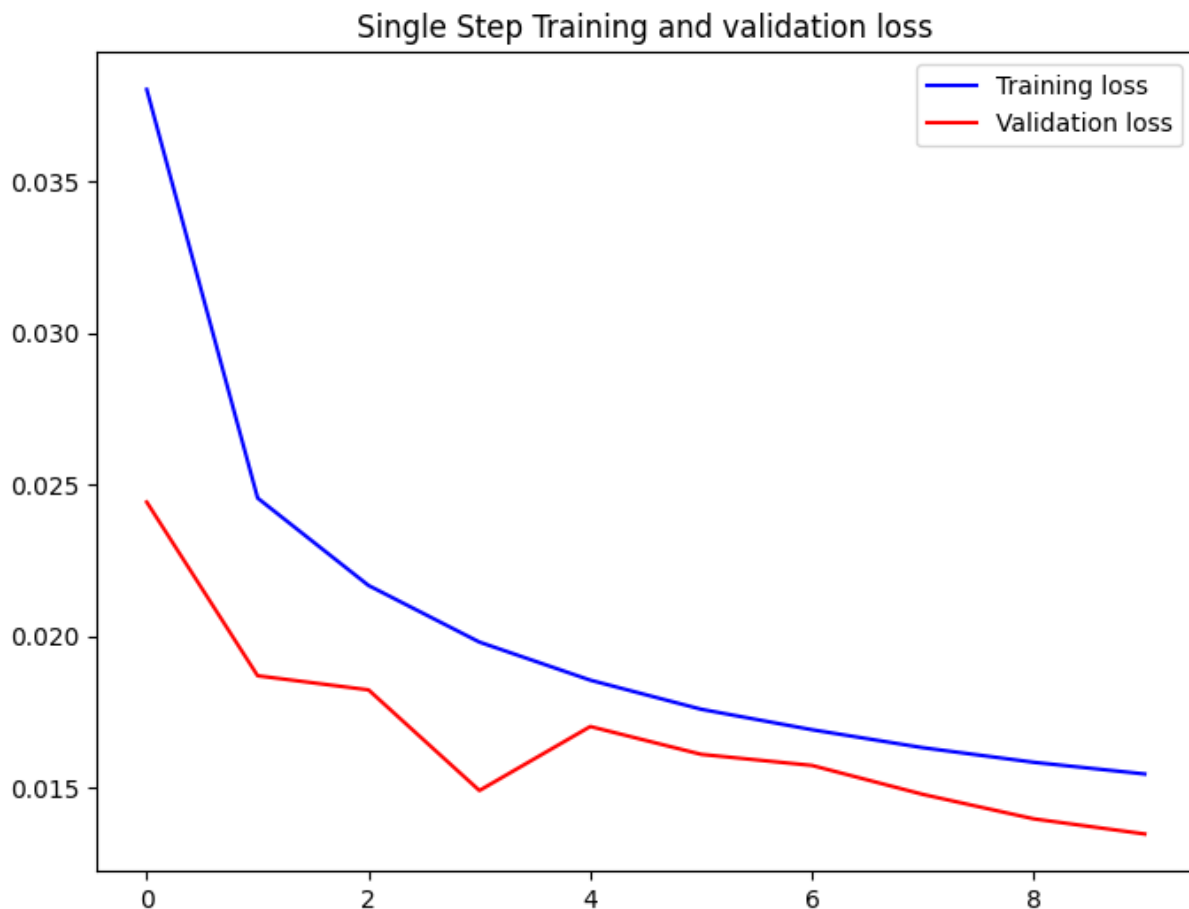
2024-05-30 09:01:30.823578: W tensorflow/core/framework/local\_rendezvous.cc:404] Local rendezvous is aborting with status: OUT\_OF\_RANGE: End of sequence

```
In [ ]: single_step_history = single_step_model.fit(train_data_single, epochs=EPOCHS,
                                                    steps_per_epoch=EVALUATION_INTERVAL,
                                                    validation_data=val_data_single,
                                                    validation_steps=50)
```

```
Epoch 1/10
2000/2000 ————— 8s 3ms/step - loss: 0.0670 - val_loss: 0.0244
Epoch 2/10
2000/2000 ————— 7s 3ms/step - loss: 0.0256 - val_loss: 0.0187
Epoch 3/10
2000/2000 ————— 7s 3ms/step - loss: 0.0222 - val_loss: 0.0182
Epoch 4/10
2000/2000 ————— 7s 3ms/step - loss: 0.0202 - val_loss: 0.0149
Epoch 5/10
2000/2000 ————— 7s 3ms/step - loss: 0.0188 - val_loss: 0.0170
Epoch 6/10
2000/2000 ————— 7s 3ms/step - loss: 0.0177 - val_loss: 0.0161
Epoch 7/10
2000/2000 ————— 7s 3ms/step - loss: 0.0170 - val_loss: 0.0157
Epoch 8/10
2000/2000 ————— 7s 3ms/step - loss: 0.0164 - val_loss: 0.0148
Epoch 9/10
2000/2000 ————— 7s 3ms/step - loss: 0.0159 - val_loss: 0.0140
Epoch 10/10
2000/2000 ————— 7s 3ms/step - loss: 0.0155 - val_loss: 0.0135
```

```
In [ ]: def plot_train_history(history, title):
        loss = history.history['loss']
        val_loss = history.history['val_loss']
        epochs = range(len(loss))
        plt.figure()
        plt.plot(epochs, loss, 'b', label='Training loss')
        plt.plot(epochs, val_loss, 'r', label='Validation loss')
        plt.title(title)
        plt.legend()
        plt.show()
```

```
In [ ]: plot_train_history(single_step_history, 'Single Step Training and validation loss')
```

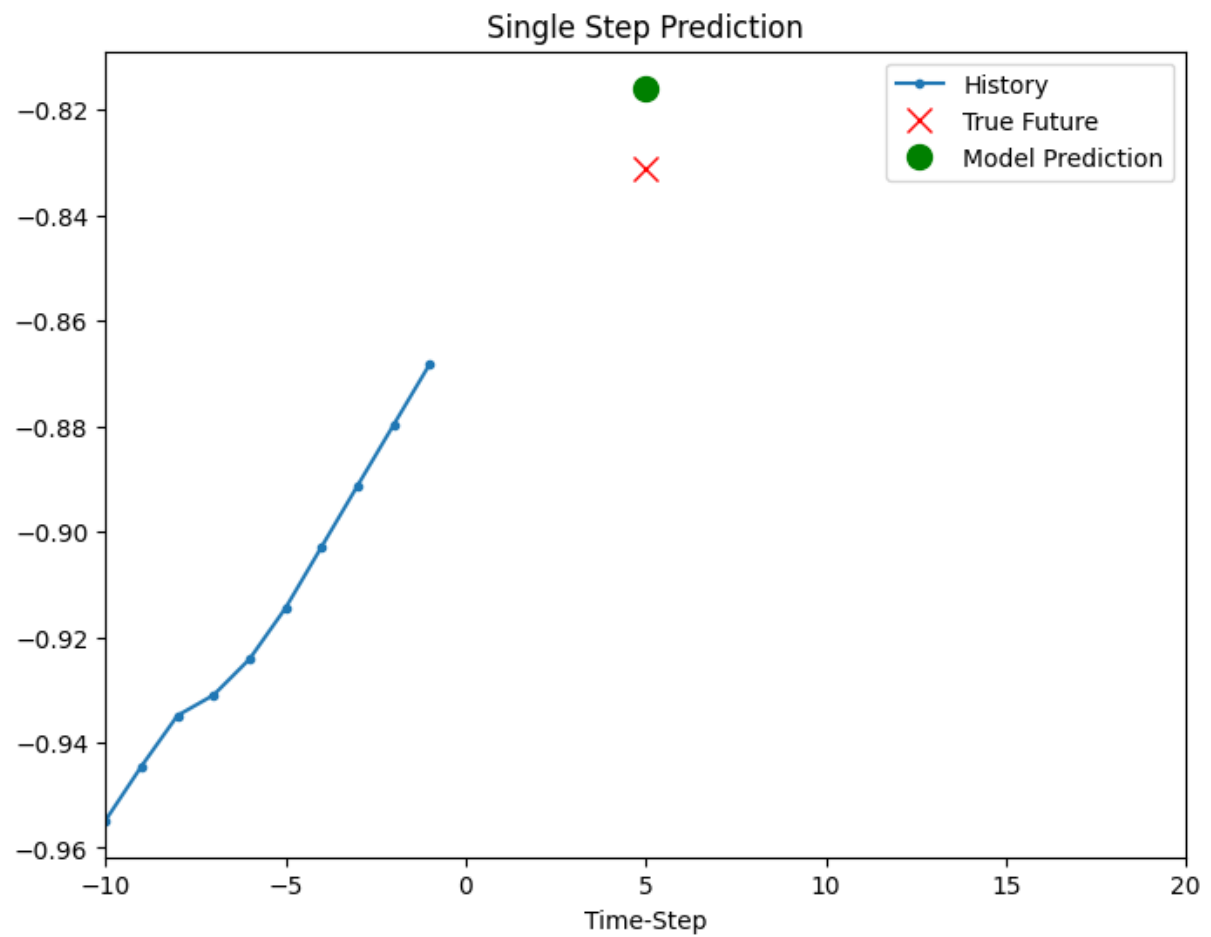


Here, we see once again a big decrease in training loss at the very beginning. However, here, the validation loss tend to shrink next to the training loss, which could indicate that the resulting network will not suffer from overfitting and is, thus, more general compared to the previous one. Moreover, I think the training needs some more epochs in order to let the loss converge at a certain value. This was probably not reached after 10 epochs.

According to time, once again it was a little bit fast than the training of the model from the original notebook, which could be due to the architecture coming with fewer parameters but also to the different machine on that the training was ran.

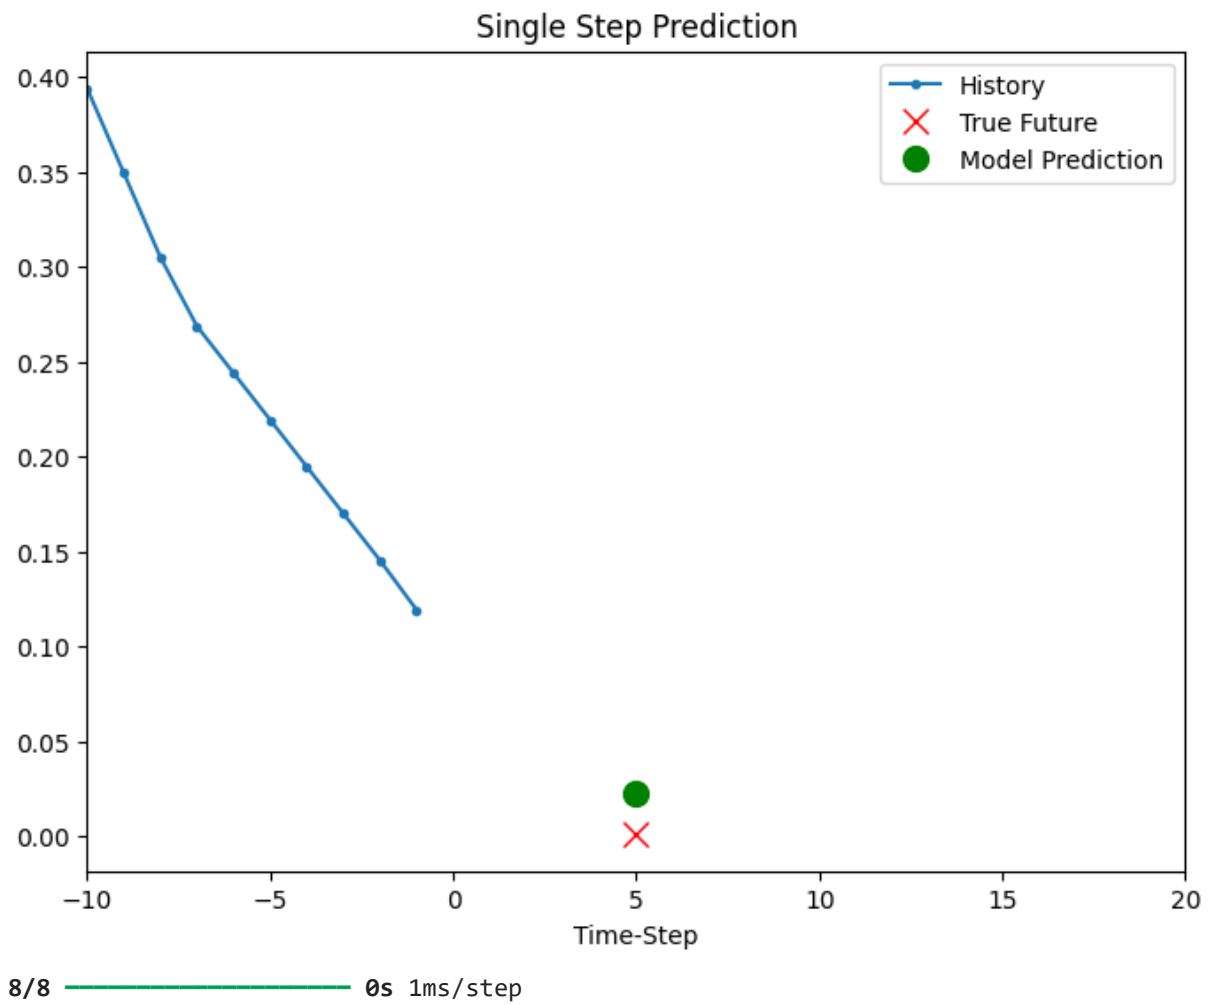
```
In [ ]: for x, y in val_data_single.take(3):  
        plot = show_plot([x[0][:, 0].numpy(), y[0].numpy(),  
                          single_step_model.predict(x)[0]], future_target,  
                          'Single Step Prediction')  
        plot.show()
```

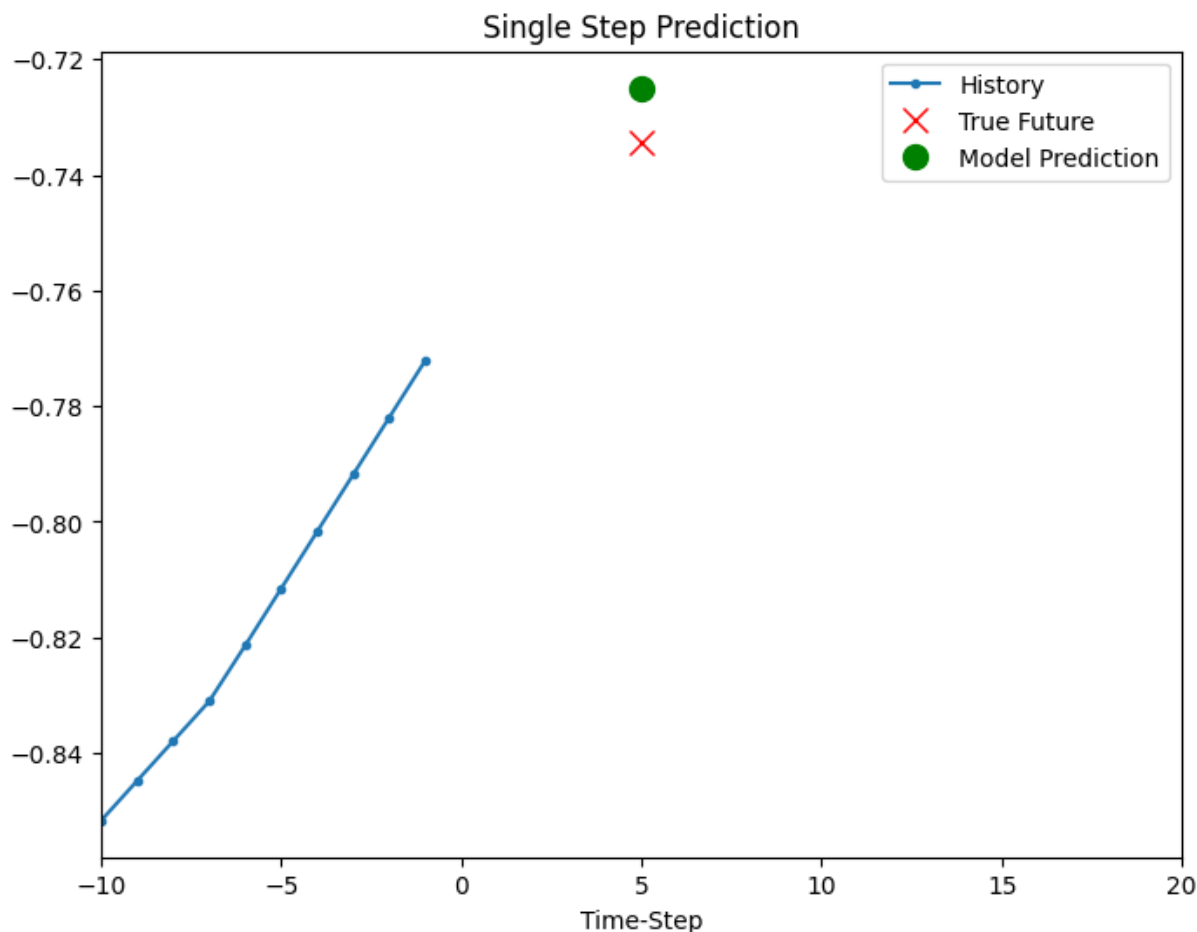
8/8 ————— 0s 1ms/step



8/8 ————— 0s 1ms/step







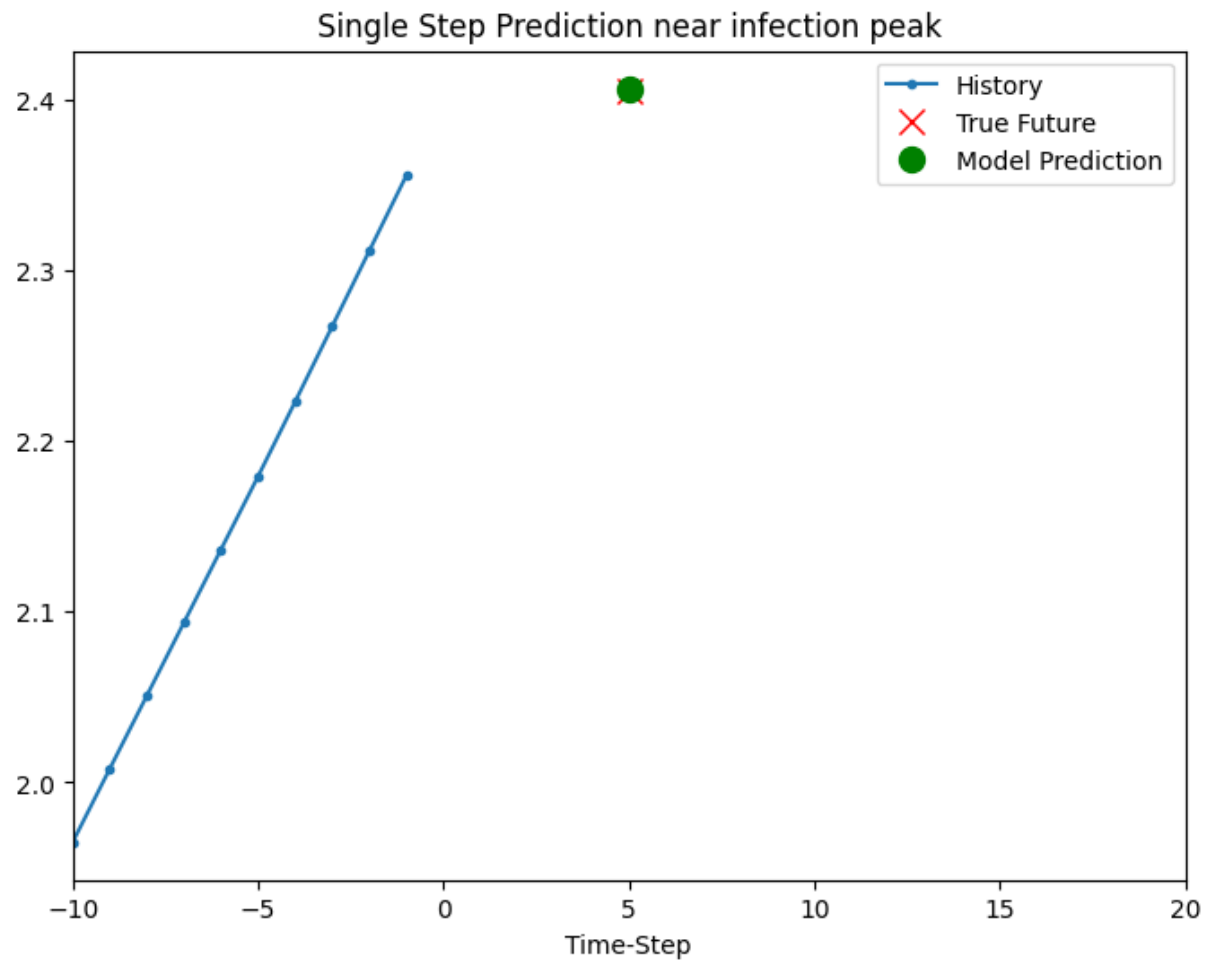
2024-05-30 09:02:40.198774: W tensorflow/core/framework/local\_rendevvous.cc:404] Local rendezvous is aborting with status: OUT\_OF\_RANGE: End of sequence

*Looking at the predictions they are a bit worse compared to the previous model. Here, we have to keep in mind, that both models performed different predictions. Here, we wanted to predict the 5 day future, while with the first model we wanted to predict the 1 day future. So maybe the error, we have seen in the previous model is likely to add up when the number of days we want to forecast into the future increases. Anyway, that the prediction is worse than in the previous model, could have been noticed when comparing the validation loss of the models.*

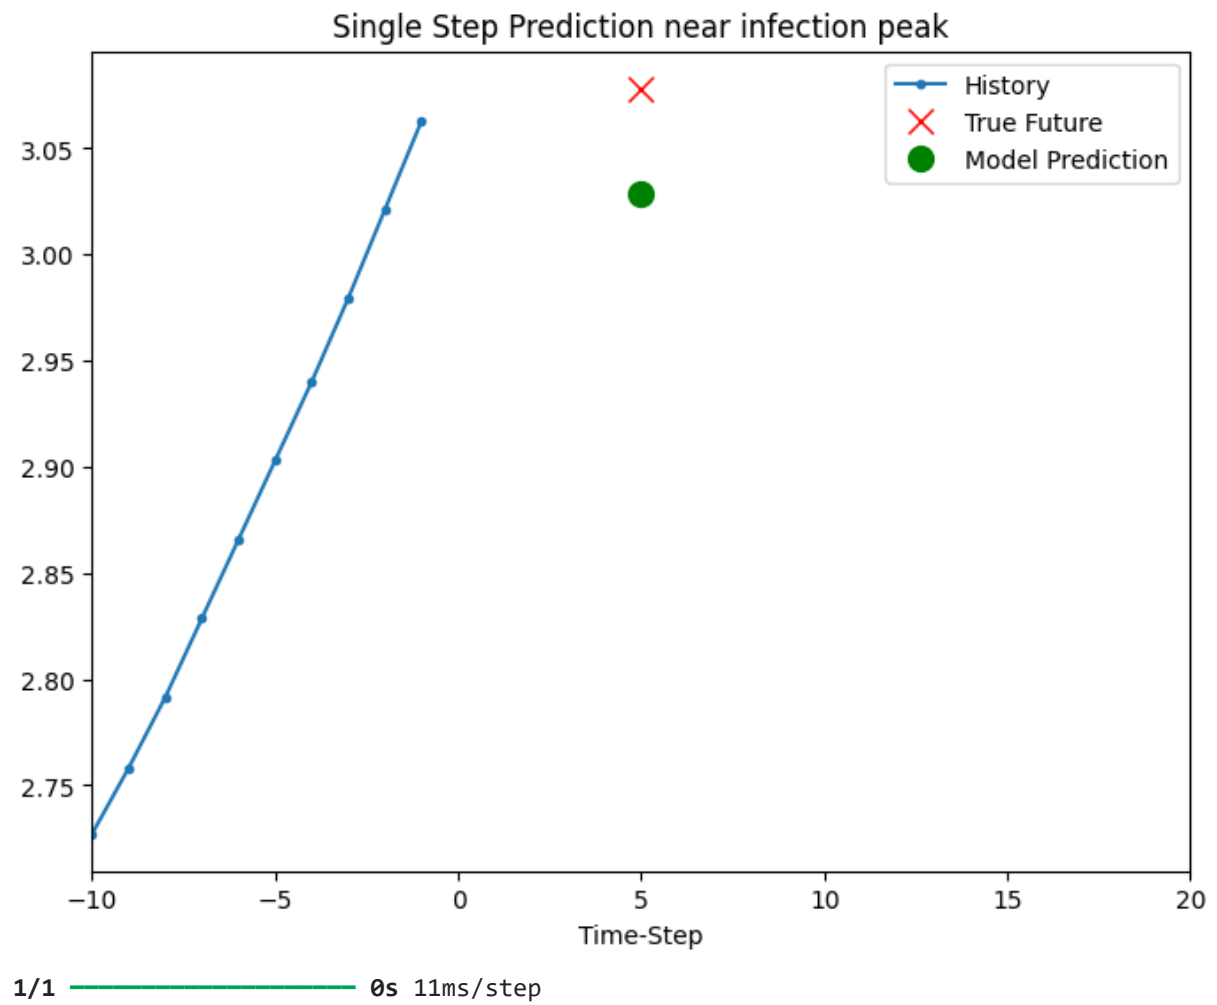
*Now, let's plot some predictions near the infection peak.*

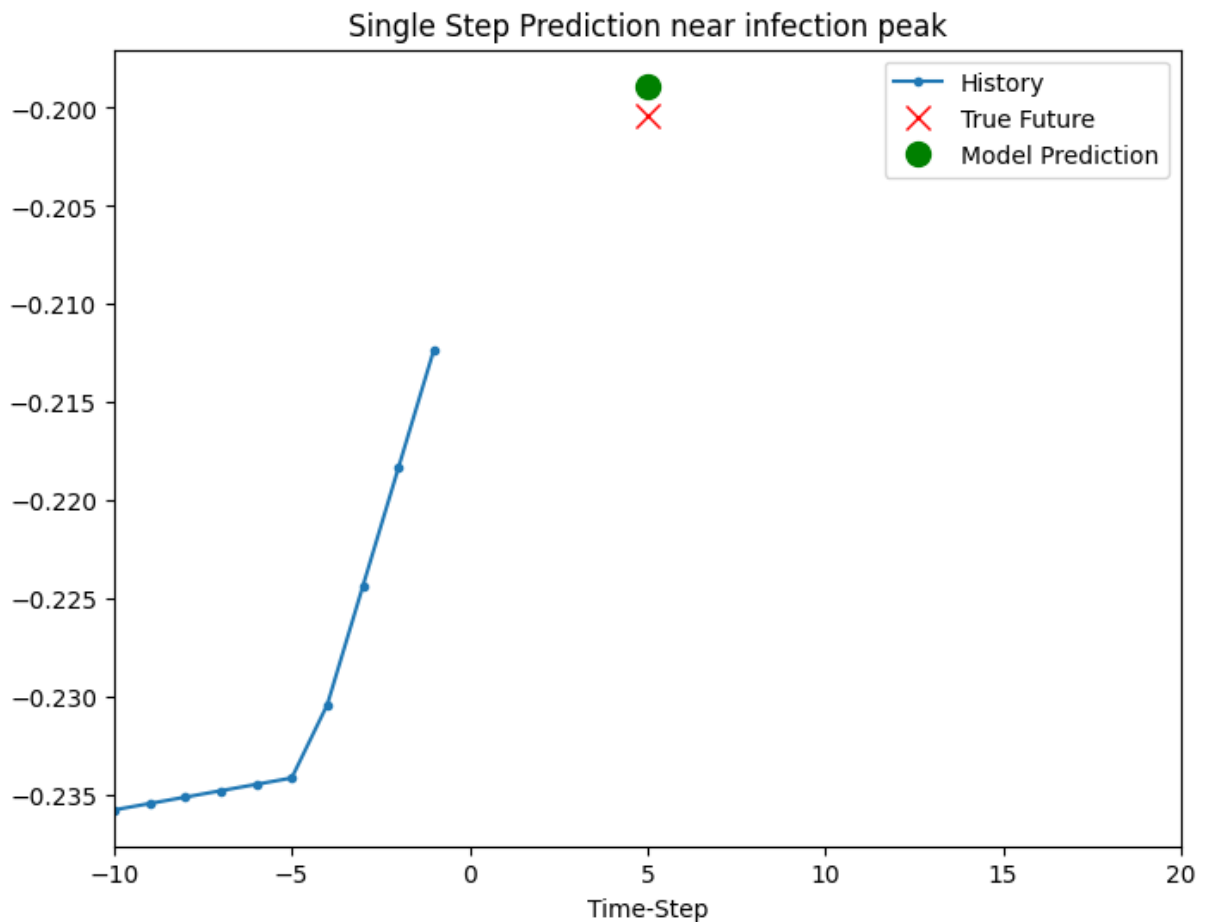
```
In [ ]: for x, y in peak_data_single.take(3):
        plot = show_plot([x[0][:, 0].numpy(), y[0].numpy(),
                          single_step_model.predict(x)[0]], future_target,
                          'Single Step Prediction near infection peak')
        plot.show()
```

1/1 ————— 0s 96ms/step



1/1 0s 13ms/step





2024-05-30 09:02:40.919216: W tensorflow/core/framework/local\_rendezvous.cc:404] Local rendezvous is aborting with status: OUT\_OF\_RANGE: End of sequence

*Here, we should take into account that the prediction made is for 3 days after the peak infection. So it is likely that the ground truth is declining in the time stamp to be predicted. Therefore, this model seems to encounter the same problem with making predictions after the direction of the trend has changed, since the second predictions is a way off compared to the true value. Maybe example 1 and 3 seems to be better, because there is some kind of plateau in infection.*

## Multivariate GRU - Multiple Steps

Still, we use a series of observed values of the three variables "Infected", "Recovered", and "Deceased" ( `past_history = 40`, `STEP = 2` ), but now to forecast the "Infected" values for a series day in the future ( `future_target = 10` ).

```
In [ ]: past_history = 40
future_target = 10
STEP = 2
x_train_multi, y_train_multi = multivariate_data(dataset, dfInfected_data, 0, TRAIN
past_history, future_target, ST
x_val_multi, y_val_multi = multivariate_data(dataset, dfInfected_data, TRAIN_SPLIT,
past_history, future_target, STEP)
```

```
In [ ]: print ('Single window of past history : {}'.format(x_train_multi[0].shape))
        print ('\nTarget window to predict : {}'.format(y_train_multi[0].shape))
        print ('\nNumber of traing data points: {}'.format(x_train_multi.shape[0]))
        print ('\nNumber of test data points: {}'.format(x_val_multi.shape[0]))
```

Single window of past history : (20, 3)

Target window to predict : (10,)

Number of traing data points: 60885

Number of test data points: 6765

*As before, retrieve data points near infection peak.*

```
In [ ]: peak_data = multivariate_data_near_peak_infection(dataset, past_history, future_tar
        peak_data.shape
```

Out[ ]: (3, 150, 51)

```
In [ ]: x_peak_multi, y_peak_multi = multivariate_data(peak_data, peak_data[0], 0, peak_dat
        past_history, future_target, STEP)
```

As before, batching and resampling; the dataset is repeated indefinitely.

```
In [ ]: train_data_multi = tf.data.Dataset.from_tensor_slices((x_train_multi, y_train_multi)
        train_data_multi = train_data_multi.cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE).

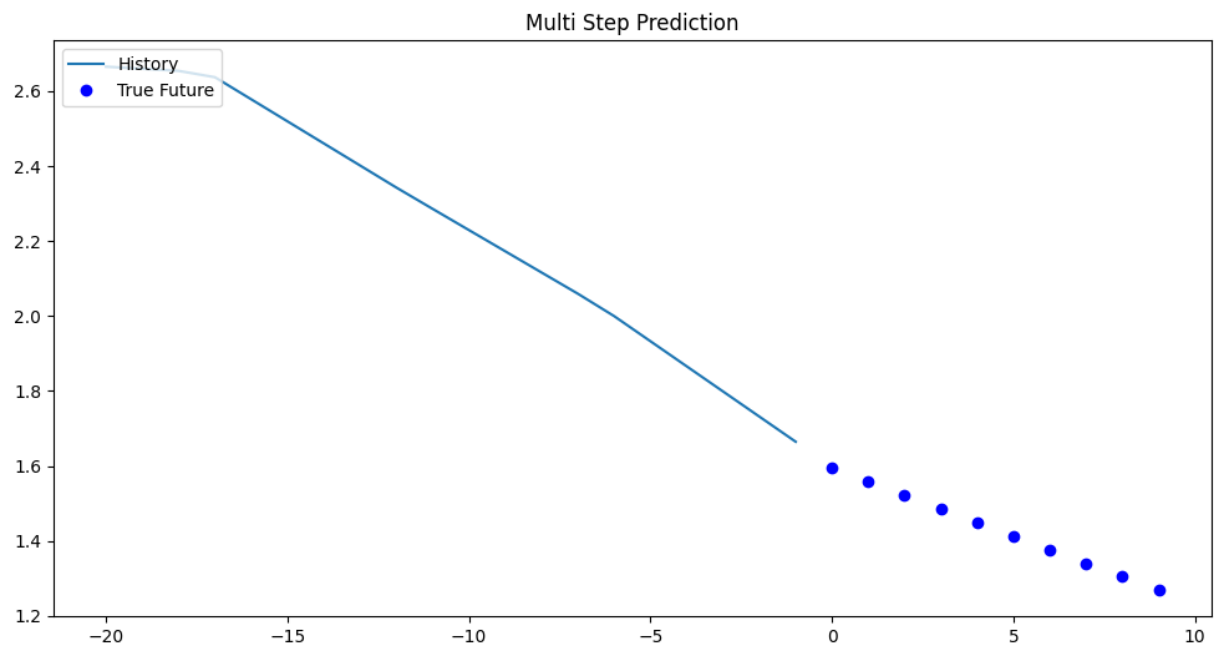
        val_data_multi = tf.data.Dataset.from_tensor_slices((x_val_multi, y_val_multi))
        val_data_multi = val_data_multi.batch(BATCH_SIZE).repeat()

        peak_data_multi = tf.data.Dataset.from_tensor_slices((x_peak_multi, y_peak_multi))
        peak_data_multi = peak_data_multi.batch(1).repeat()
```

```
In [ ]: def multi_step_plot(history, true_future, prediction, title='Multi Step Prediction')
        plt.figure(figsize=(12, 6))
        num_in = create_time_steps(len(history))
        num_out = len(true_future)
        plt.plot(num_in, np.array(history[:, 0]), label='History')
        plt.plot(np.arange(num_out), np.array(true_future), 'bo', label='True Future')
        if prediction.any():
            plt.plot(np.arange(num_out), np.array(prediction), 'ro', label='Predicted F
        plt.legend(loc='upper left')
        plt.title(title)
        plt.show()
```

```
In [ ]: for x, y in train_data_multi.take(1):
        multi_step_plot(x[0], y[0], np.array([0]))
```

```
2024-05-30 09:02:41.757811: W tensorflow/core/kernels/data/cache_dataset_ops.cc:858]
The calling iterator did not fully read the dataset being cached. In order to avoid
unexpected truncation of the dataset, the partially cached contents of the dataset
will be discarded. This can happen if you have an input pipeline similar to `dataset
t.cache().take(k).repeat()`. You should use `dataset.take(k).cache().repeat()` inste
ad.
```



```
2024-05-30 09:02:41.847356: W tensorflow/core/framework/local_rendezvous.cc:404] Loc
al rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
```

Now we build a model with two *GRU* layers.

```
In [ ]: multi_step_model = tf.keras.models.Sequential()
multi_step_model.add(tf.keras.layers.GRU(32,
                                          return_sequences=True,
                                          input_shape=x_train_multi.shape[-2:]))
multi_step_model.add(tf.keras.layers.GRU(16, activation='relu'))
multi_step_model.add(tf.keras.layers.Dense(future_target))

multi_step_model.compile(optimizer=tf.keras.optimizers.RMSprop(clipvalue=1.0), loss
multi_step_model.summary()
x_train_multi.shape[-2:]
```

Model: "sequential\_11"

Layer (type)	Output Shape	Param #
gru_14 (GRU)	(None, 20, 32)	3,552
gru_15 (GRU)	(None, 16)	2,400
dense_11 (Dense)	(None, 10)	170

Total params: 6,122 (23.91 KB)

Trainable params: 6,122 (23.91 KB)

Non-trainable params: 0 (0.00 B)

```
Out[ ]: (20, 3)
```

*As before, we observe that the number of parameter is smaller compared to the model with the LSTM layers, due to the different architectures.*

```
In [ ]: for x, y in val_data_multi.take(1):  
        print (multi_step_model.predict(x).shape)
```

```
8/8 ————— 0s 2ms/step  
(256, 10)
```

```
2024-05-30 09:02:42.115515: W tensorflow/core/framework/local_rendevvous.cc:404] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
```

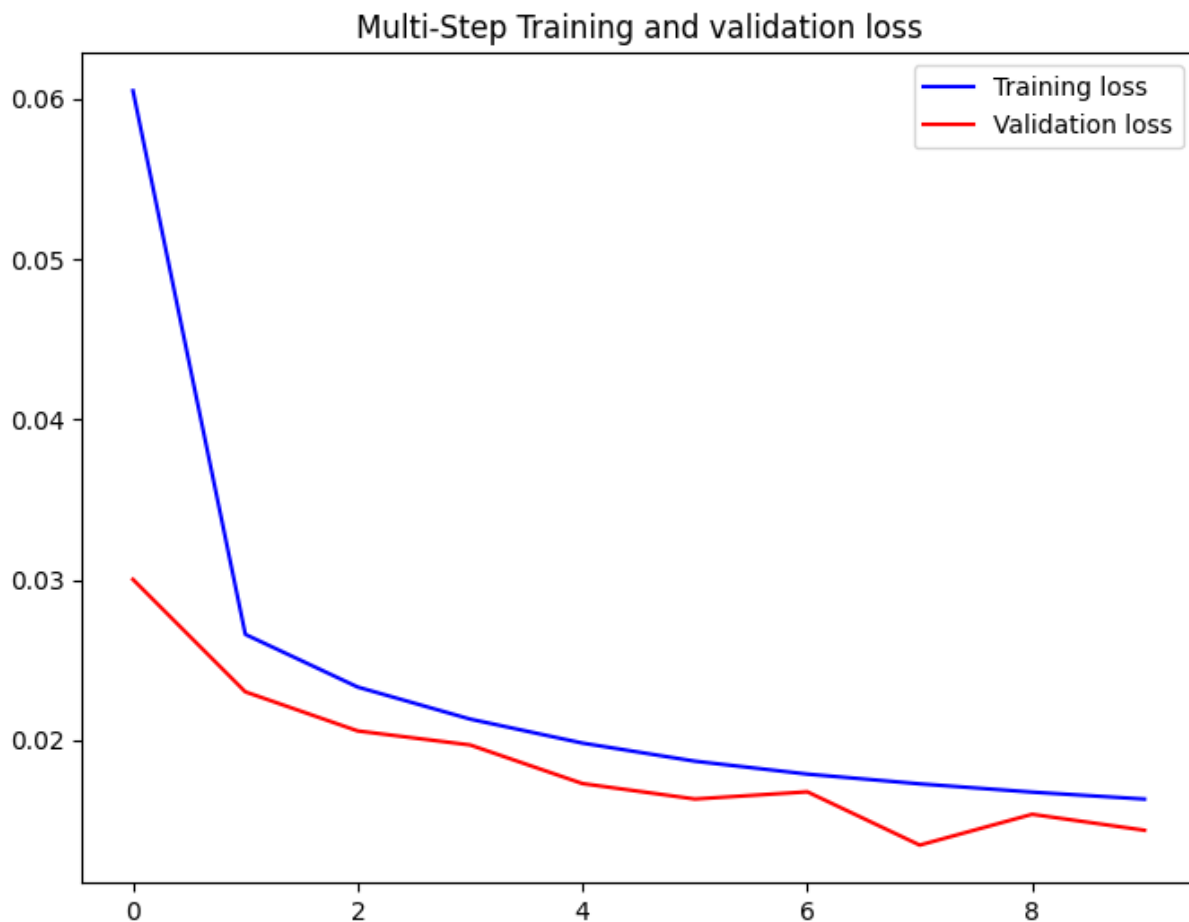
The training time is longer for this more complex model.

```
In [ ]: multi_step_history = multi_step_model.fit(train_data_multi, epochs=EPOCHS,  
                                                steps_per_epoch=EVALUATION_INTERVAL,  
                                                validation_data=val_data_multi,  
                                                validation_steps=50)
```

```
Epoch 1/10  
2000/2000 ————— 22s 10ms/step - loss: 0.1428 - val_loss: 0.0300  
Epoch 2/10  
2000/2000 ————— 20s 10ms/step - loss: 0.0277 - val_loss: 0.0230  
Epoch 3/10  
2000/2000 ————— 20s 10ms/step - loss: 0.0239 - val_loss: 0.0206  
Epoch 4/10  
2000/2000 ————— 21s 10ms/step - loss: 0.0217 - val_loss: 0.0197  
Epoch 5/10  
2000/2000 ————— 20s 10ms/step - loss: 0.0202 - val_loss: 0.0173  
Epoch 6/10  
2000/2000 ————— 20s 10ms/step - loss: 0.0190 - val_loss: 0.0163  
Epoch 7/10  
2000/2000 ————— 20s 10ms/step - loss: 0.0181 - val_loss: 0.0168  
Epoch 8/10  
2000/2000 ————— 21s 10ms/step - loss: 0.0174 - val_loss: 0.0135  
Epoch 9/10  
2000/2000 ————— 21s 10ms/step - loss: 0.0169 - val_loss: 0.0154  
Epoch 10/10  
2000/2000 ————— 20s 10ms/step - loss: 0.0164 - val_loss: 0.0144
```

```
In [ ]: plot_train_history(multi_step_history, 'Multi-Step Training and validation loss')
```



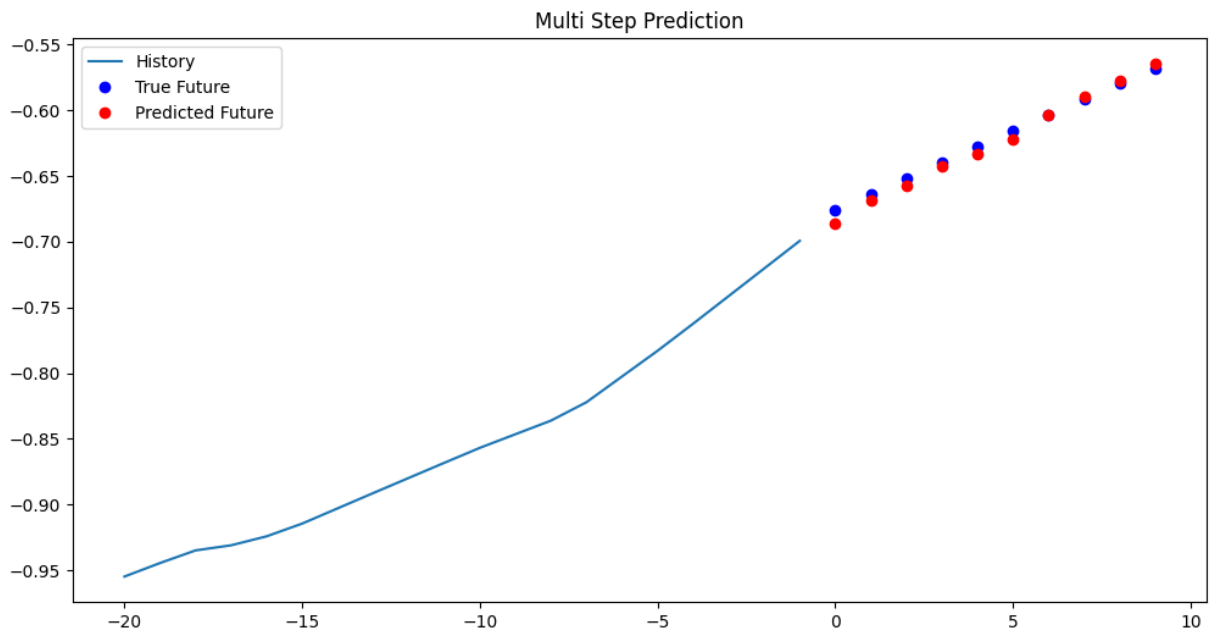


Here, we see once again a big decrease in training loss at the very beginning. The validation loss tend to shrink next to the training loss, which could indicate that the resulting network will not suffer from overfitting. Moreover, it is possible that the training needs some more epochs in order to let the loss converge at a certain value. However, the situation here looks quite better (regarding to that the model already reached to optimum ) compared to the previous model.

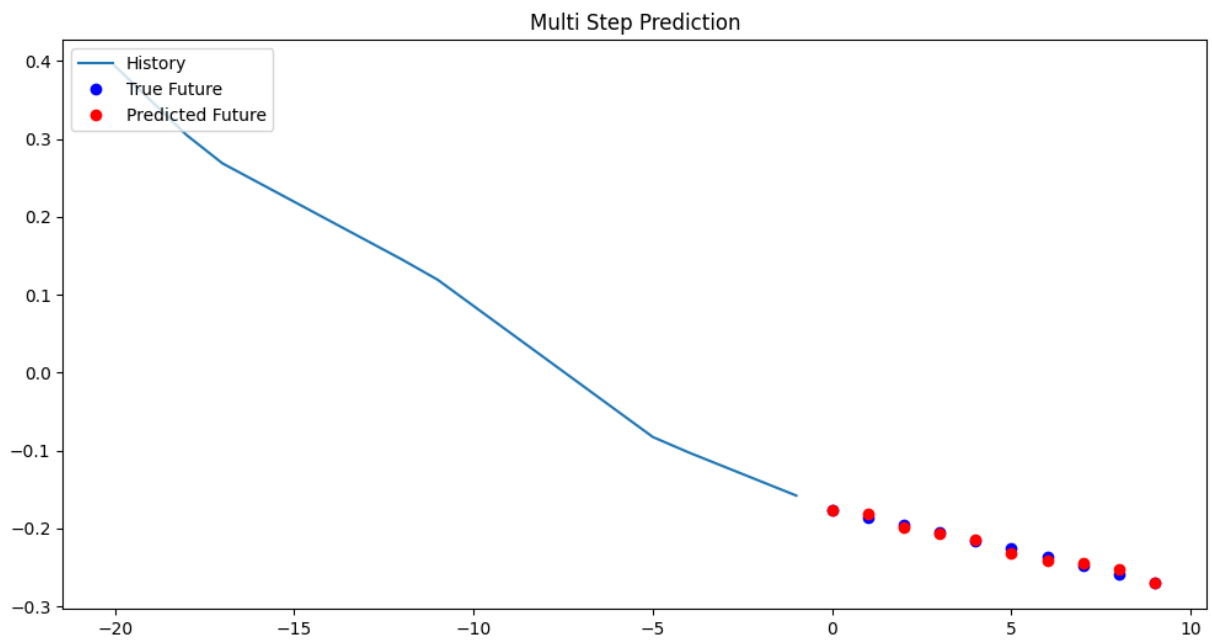
According to time, the training was slower compared to the other models. This is due to the more complex architecture which come with more parameters and , thus, a higher computational effort. Anyway, once again it was a little bit fast than the training of the model from the original notebook, which could be due to the architecture coming with fewer parameters but also to the different machine on that the training was ran.

```
In [ ]: for x, y in val_data_multi.take(3):  
        multi_step_plot(x[0], y[0], multi_step_model.predict(x)[0])
```

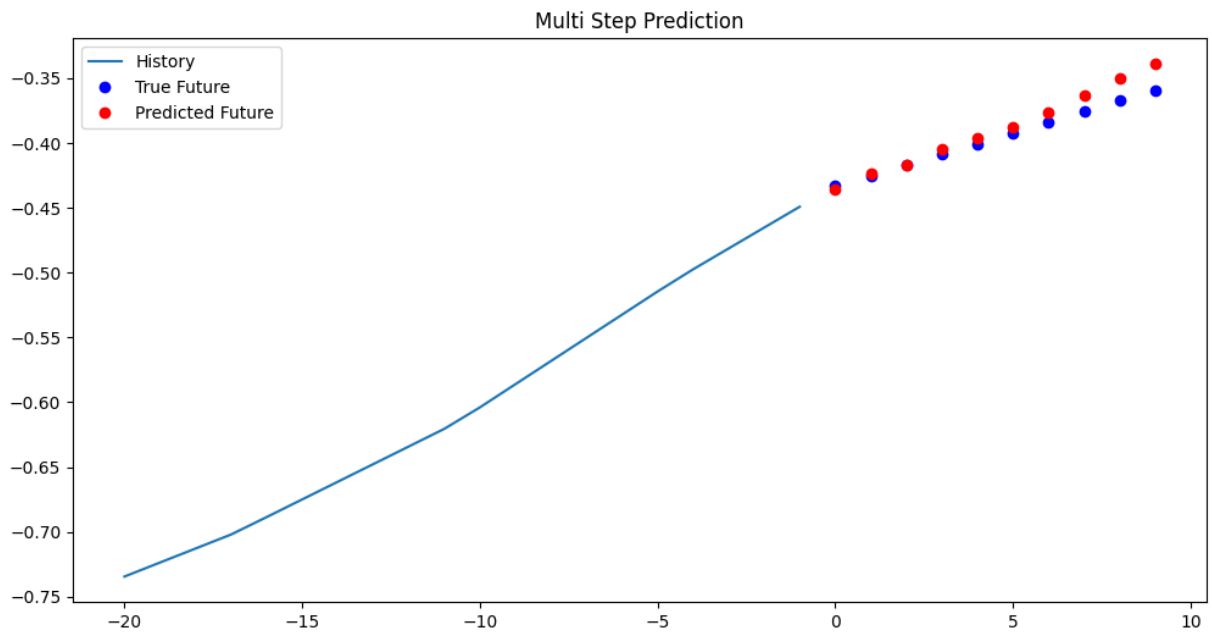
8/8 ————— 0s 2ms/step



8/8 ————— 0s 2ms/step



8/8 ————— 0s 2ms/step



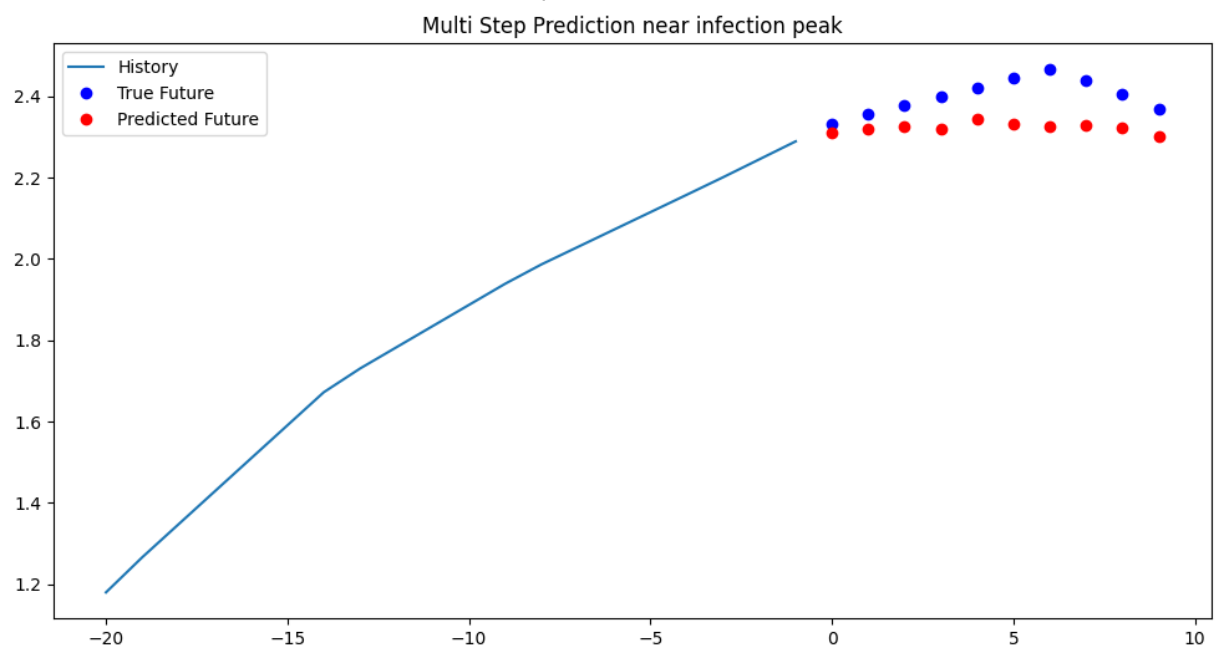
2024-05-30 09:06:08.637510: W tensorflow/core/framework/local\_rendevvous.cc:404] Local rendezvous is aborting with status: OUT\_OF\_RANGE: End of sequence

*Looking at the predictions they are quite good. In some cases, one sees that the difference between true value and prediction is larger when forecasting is made for a more distant future (third example).*

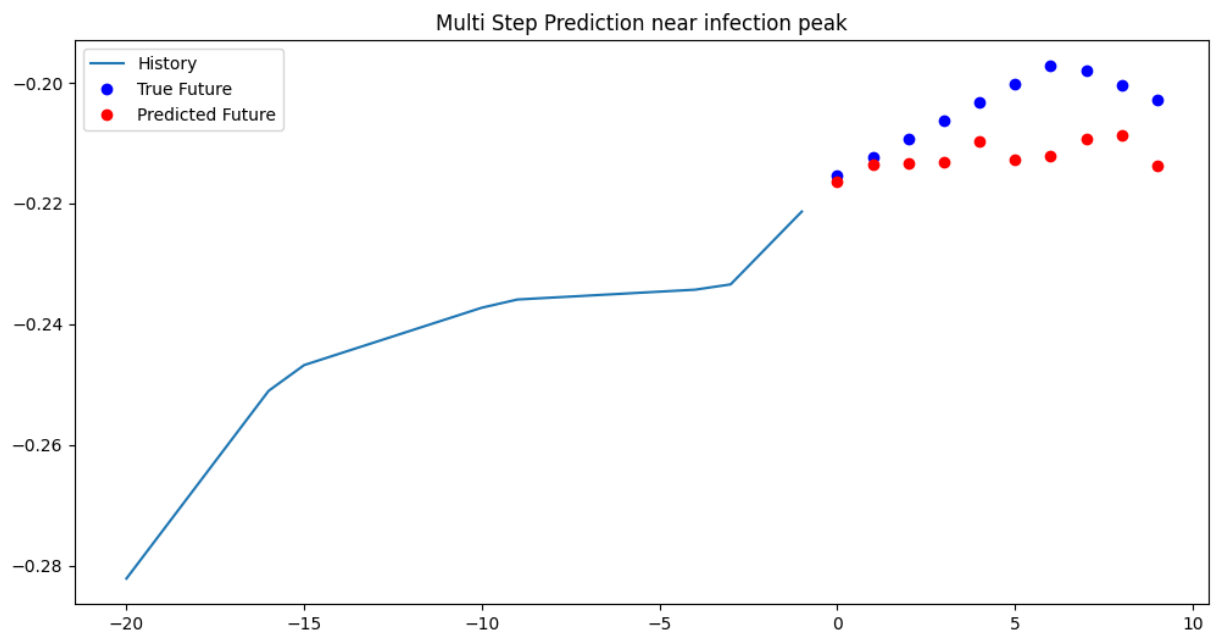
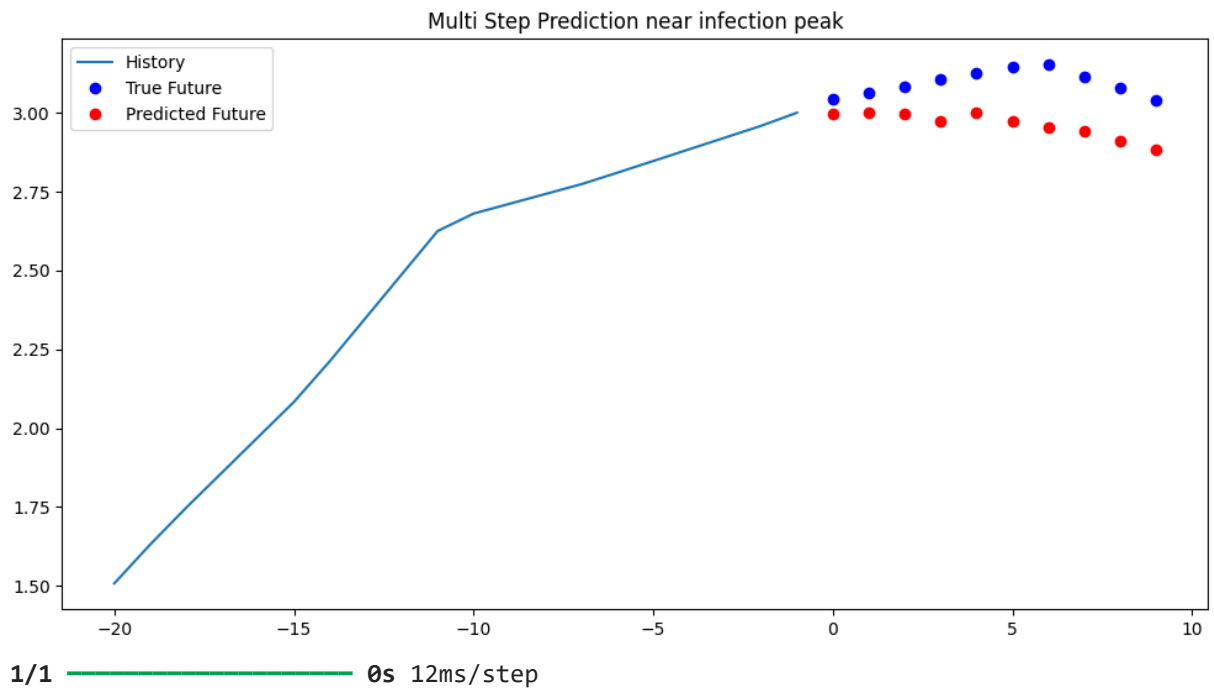
*Now, let's look at the prediction near infection peak.*

```
In [ ]: for x, y in peak_data_multi.take(3):  
        multi_step_plot(x[0], y[0], multi_step_model.predict(x)[0], title="Multi Step P
```

1/1 ————— 0s 167ms/step



1/1 ————— 0s 11ms/step



2024-05-30 09:06:09.171313: W tensorflow/core/framework/local\_rendezvous.cc:404] Local rendezvous is aborting with status: OUT\_OF\_RANGE: End of sequence

Here, we see that the model is likely to miss the peak and predicts a more flat trajectory. So once again, it seems to have trouble with the change in the direction of the trajectory, i.e. from increasing infections to decreasing infections. This could be due to the training data. As we have seen in the plotted examples, there are some timelines that are likely to haven't reached the peak yet, and, therefore, were increasing all the time. So maybe, one could try to focus the training on timelines showing a global within the, e.g., first 300 days, in order to ensure a complete infection cycle - if those data exists.